

Distributed Systems

Techniques for scaling

Replication and caching: Make copies of data available at different machines

- Replicated **file servers and databases**
- Mirrored **Web sites**
- Web caches (in **browsers** and proxies)
- **File caching** (at server and client)

Scaling: The problem with replication

Applying replication is easy, except for one thing

- Having multiple copies (cached or replicated), leads to **inconsistencies**:
- **Modifying one copy** makes that copy different from the rest.
- Always keeping copies consistent and in a general way requires global **synchronization** on each modification.
- Global synchronization precludes large-scale solutions.

Developing distributed systems

False (and often hidden) assumptions



- The network is reliable
- The network is secure
- The network is homogeneous
- The topology does not change
- Latency is zero
- Bandwidth is infinite
- Transport cost is zero
- There is one administrator

Three types of distributed systems

- High performance distributed computing systems
HDCS
- Distributed information systems
DIS
- Distributed systems for pervasive computing
DSPC

Parallel computing

- High-performance distributed computing started with parallel computing



Distributed shared memory systems

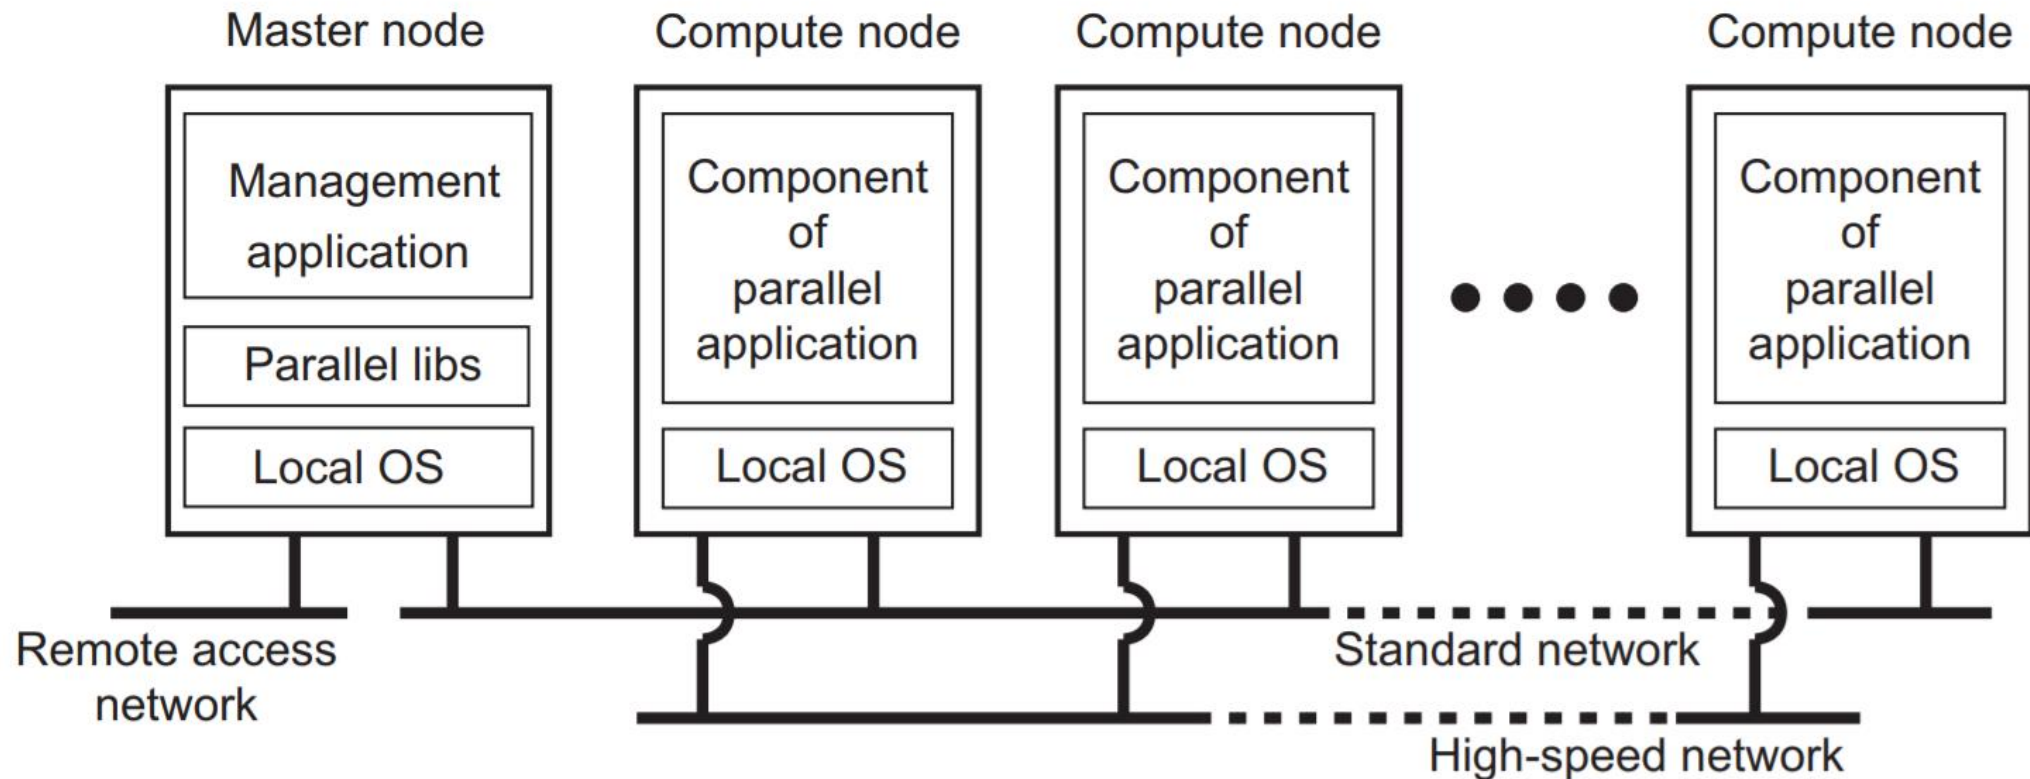
- Multiprocessors are relatively **easy to program**
- In comparison to multicomputers,
- yet have problems when **increasing the number of processors (or cores)**.
- **Solution**: Try to implement a shared-memory model on top of a multicomputer.

Problem

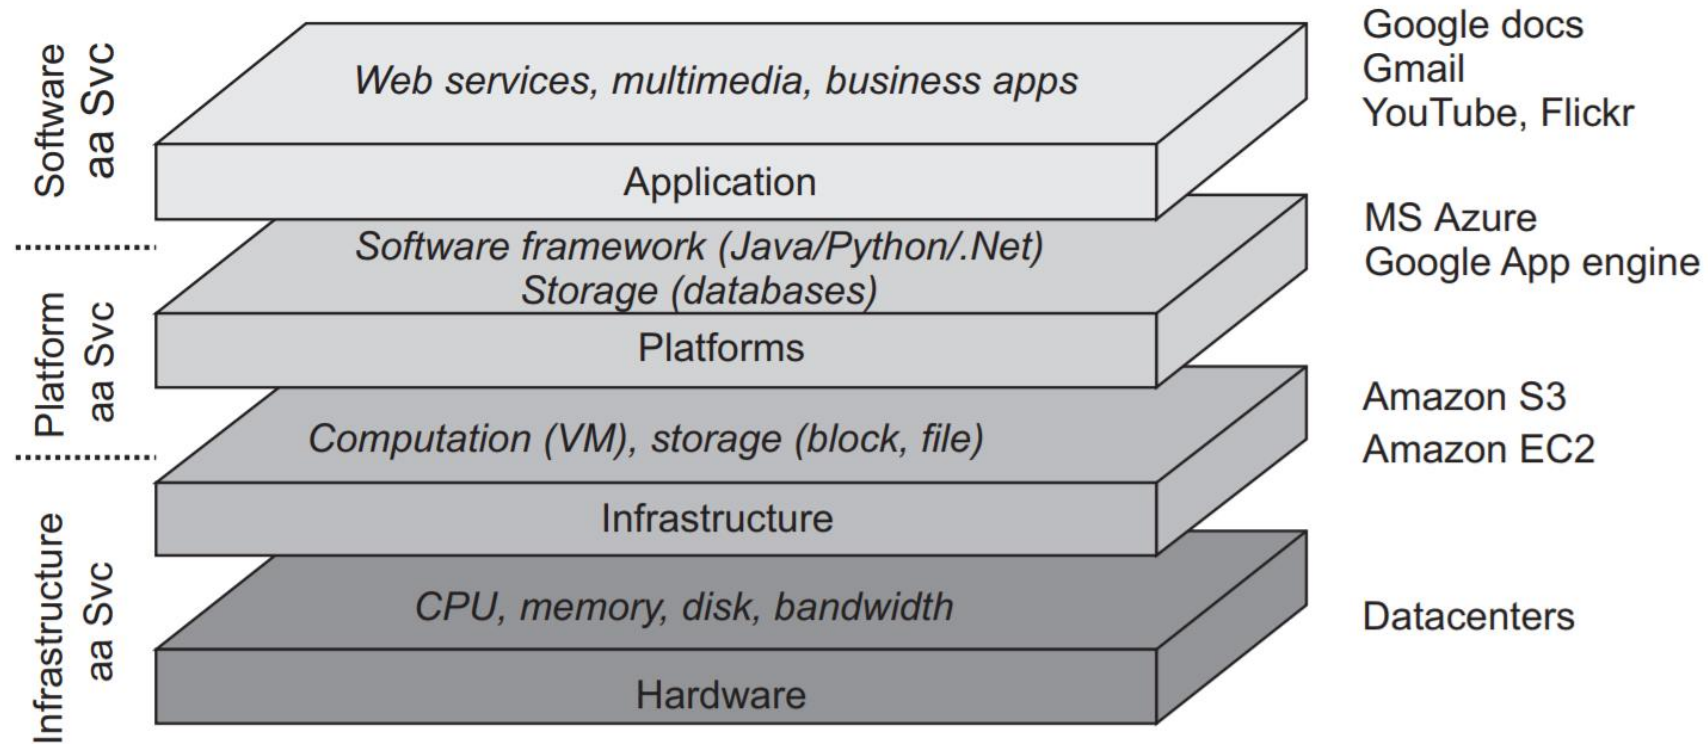
- Performance of distributed shared memory could never compete with that of **multiprocessors**, and failed to meet the expectations of programmers.
- It has been widely abandoned by now.

Cluster computing

- Homogeneous: same OS, near-identical hardware
- Single managing node



Cloud computing



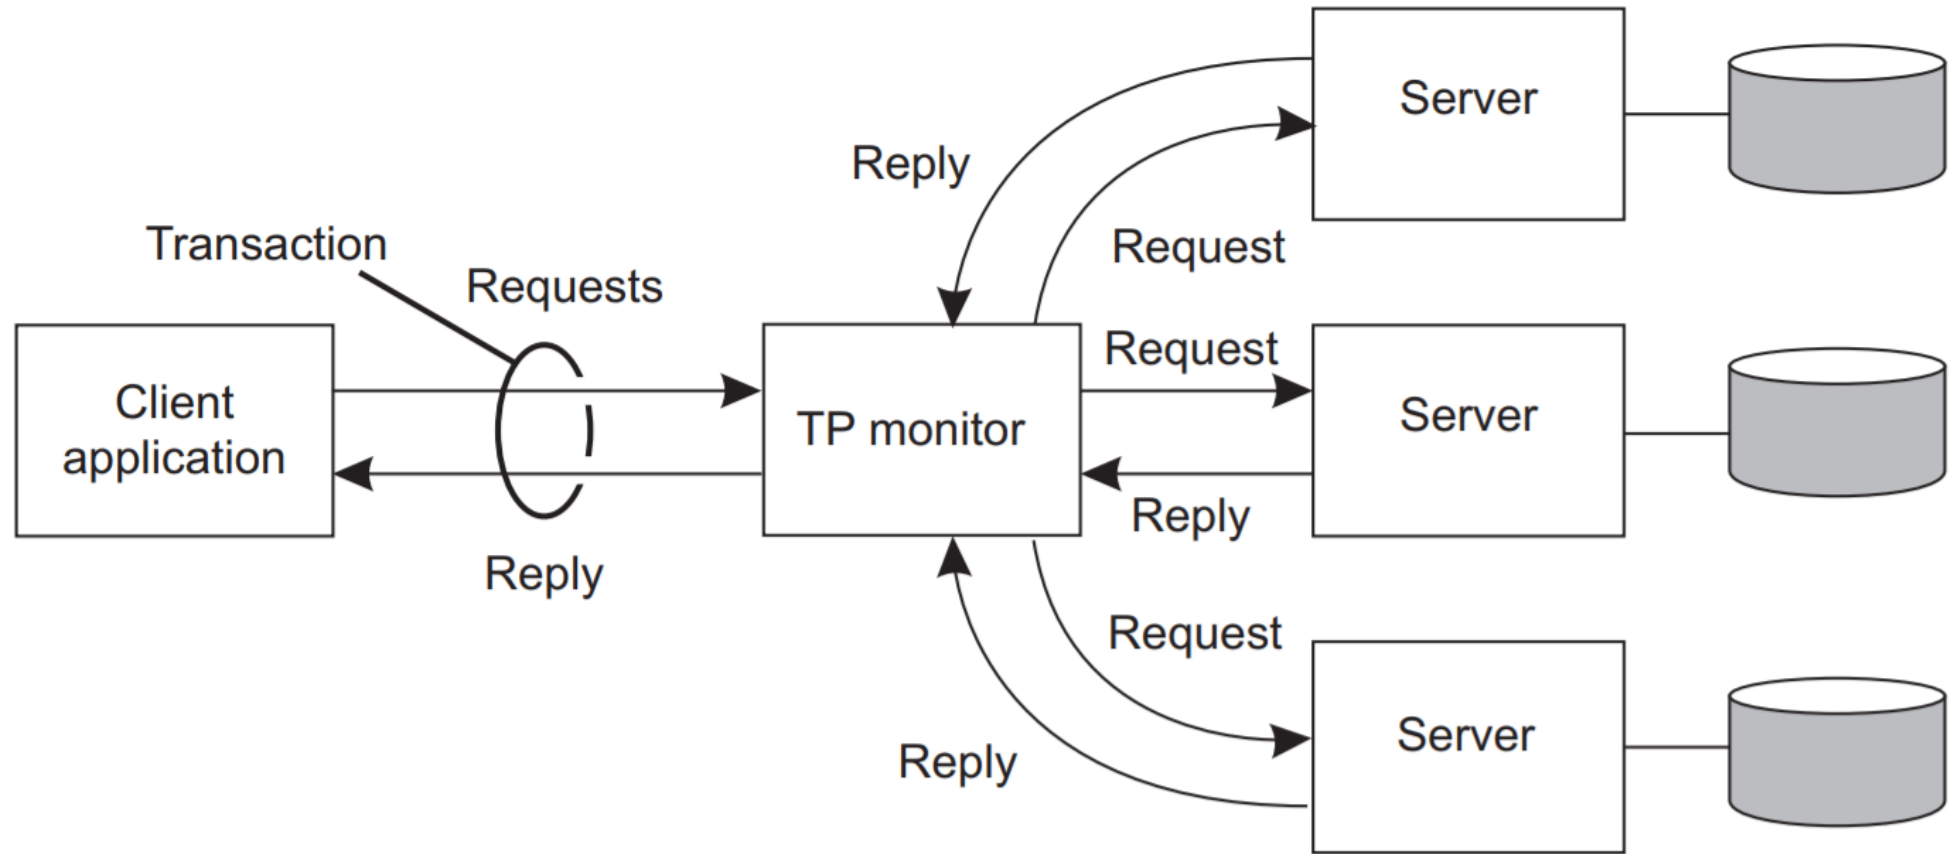
Cloud computing

- **Hardware**: Processors, routers, power and cooling systems. Customers normally never get to see these.
- **Infrastructure**: Deploys virtualization techniques. Evolves around allocating and managing virtual storage devices and virtual servers.
- **Platform**: Provides higher-level abstractions for storage and such. Example: Amazon S3 storage system offers an API for (locally created) files to be organized and stored in so-called **buckets**.
- **Application**: Actual applications, such as office suites (text processors, spreadsheet applications, presentation applications). Comparable to the suite of apps shipped with OSes.

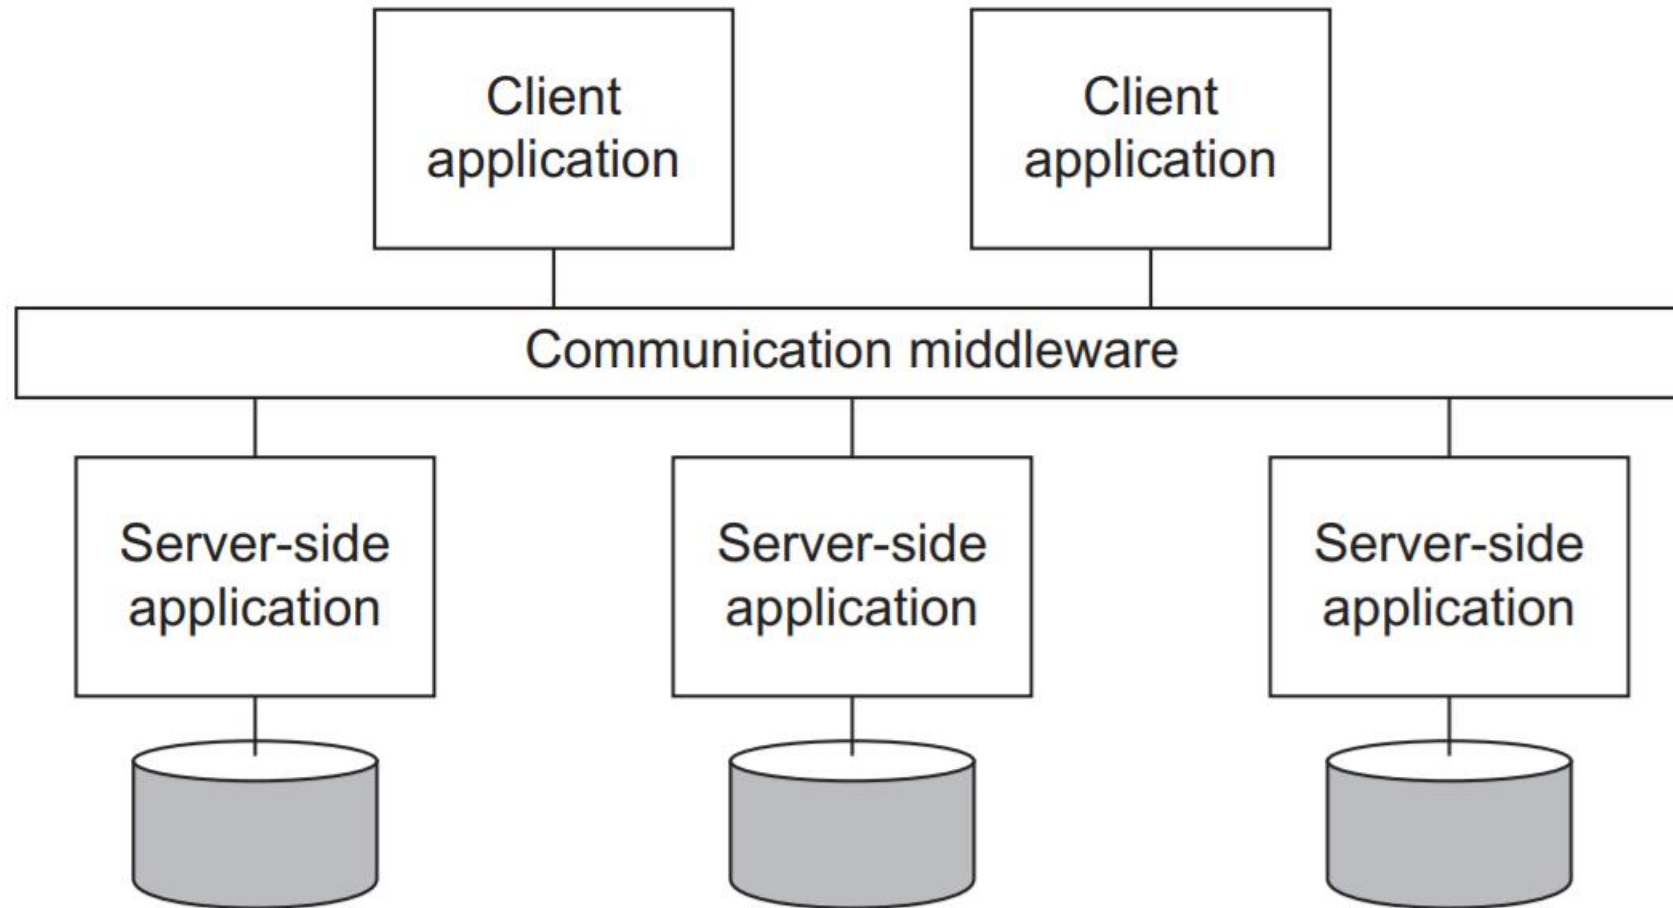
Integrating applications

- A networked application is one that runs on a server making its services available to remote clients.
- **Simple integration**: clients combine requests for (different) applications; send that off; collect responses, and present a coherent result to the user.
- Allow direct application-to-application communication, leading to Enterprise Application Integration.

TPM: Transaction Processing Monitor



Middleware



Middleware

Remote Procedure Call (RPC):

- Requests are sent through local procedure call, packaged as message, processed, responded through message, and result returned as return from call.

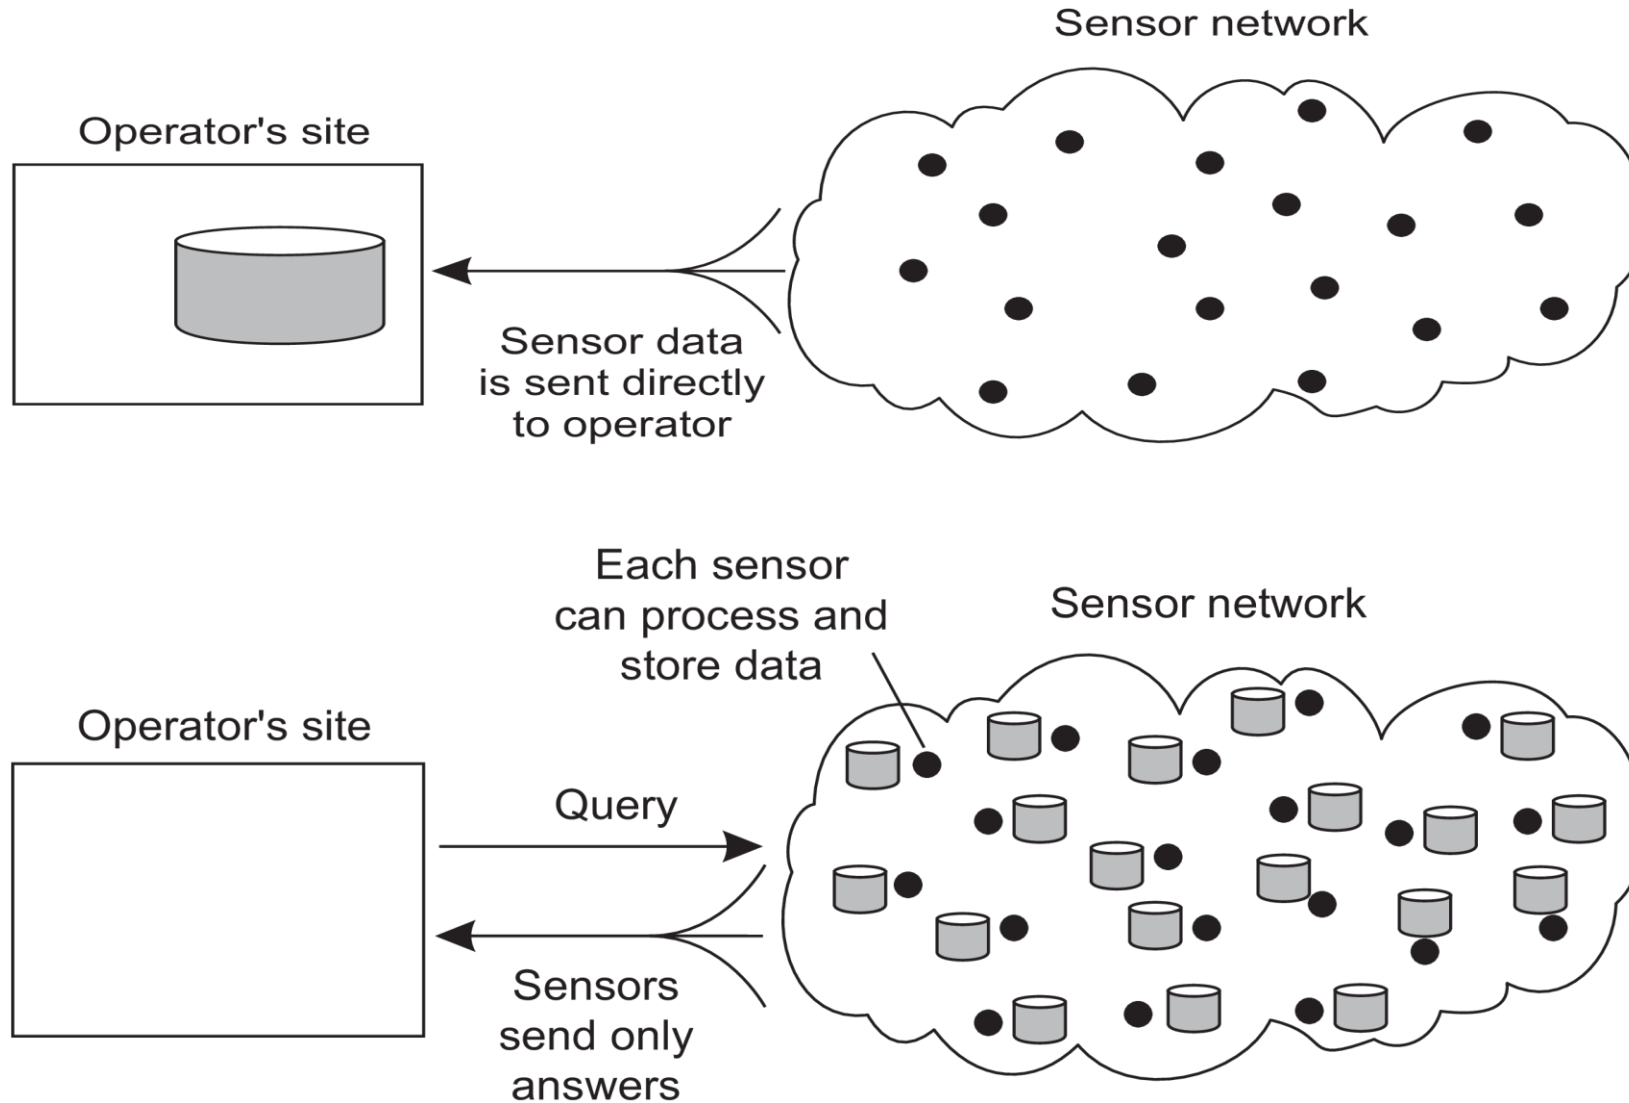
Message Oriented Middleware (MOM):

- Messages are sent to logical contact point (**published**), and forwarded to **subscribed** applications.

Distributed pervasive systems

- **Ubiquitous computing systems**: pervasive and **continuously present**, i.e., there is a continuous interaction between system and user.
- **Mobile computing systems**: pervasive, but emphasis is on the fact that devices are **inherently mobile**.
- **Sensor (and actuator) networks**: pervasive, with emphasis on the actual (collaborative) **sensing and actuation** of the environment.


Sensor networks as distributed databases



Architectural styles

- (replaceable) components with well-defined interfaces
- the way that components are connected to each other
- the data exchanged between components
- how these components and connectors are jointly configured into a system.

Example: communication protocols

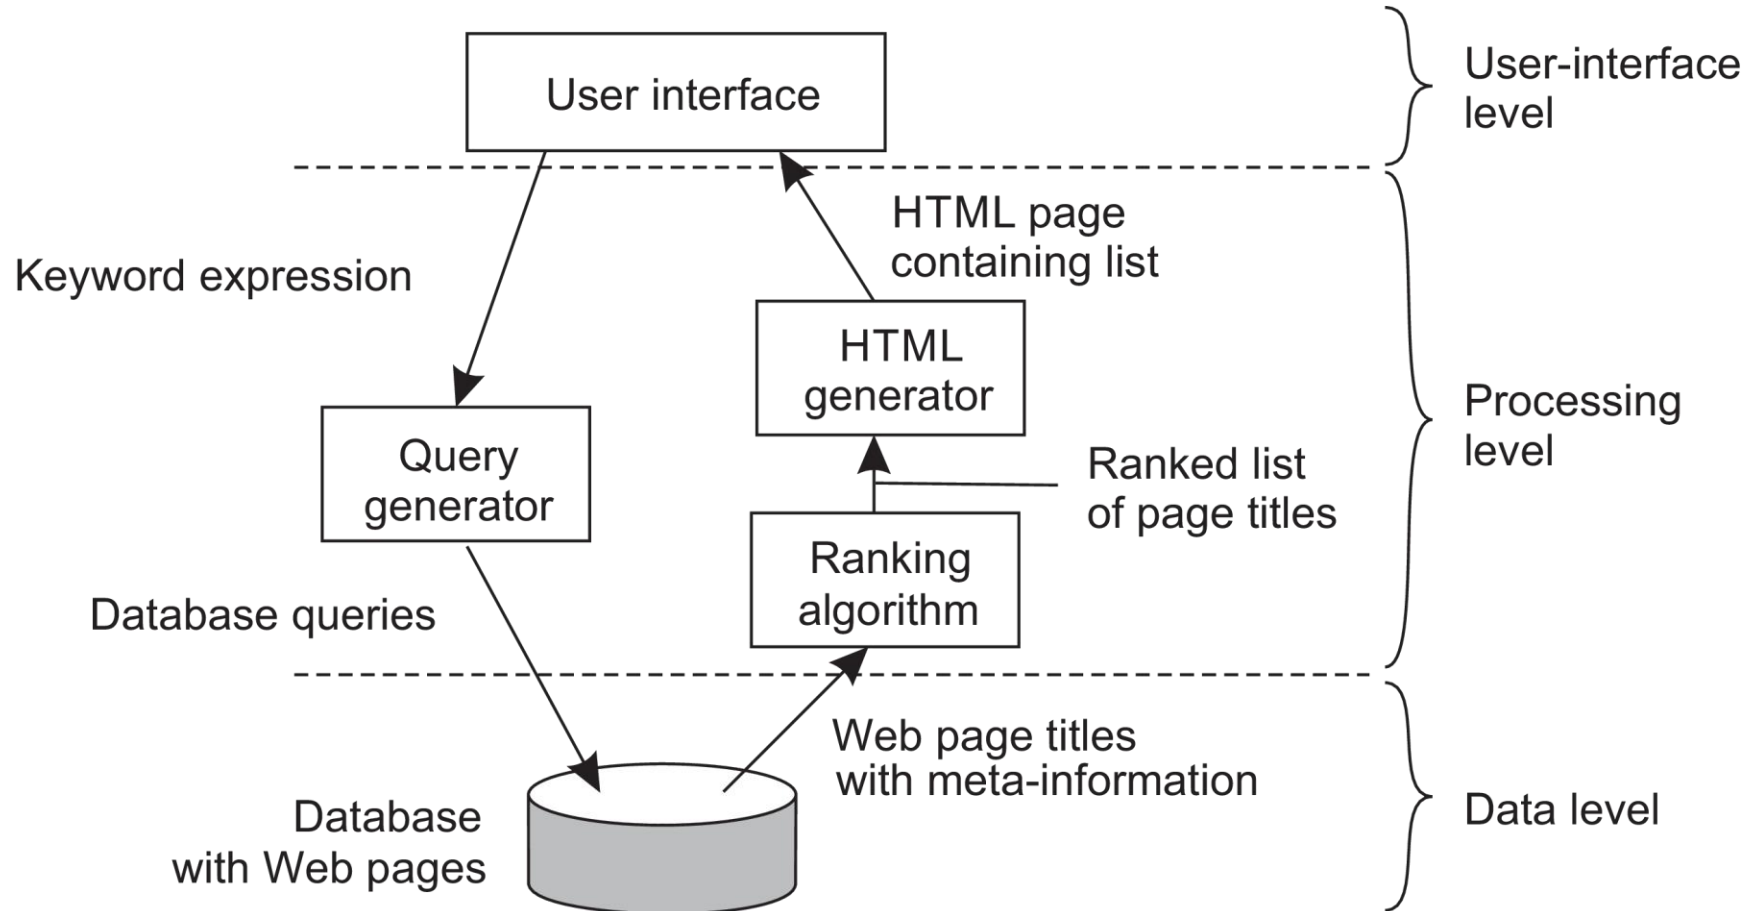


Application Layering

- **Application-interface layer** contains units for interfacing to users or external applications
- **Processing layer** contains the functions of an application, i.e., without specific data
- **Data layer** contains the data that a client wants to manipulate through the application components

This layering is found in many distributed information systems, using traditional database technology and accompanying applications.

Application Layering



RESTful architectures

- View a distributed system as a collection of resources, individually managed by components.
- Resources may be added, removed, retrieved, and modified by (remote) applications.

Operation	Description
POST	Create a new resource
GET	Retrieve the state of a resource in some representation
DELETE	Delete a resource
PUT	Modify a resource by transferring a new state