

Distributed Systems

Example: Amazon's Simple Storage Service

Typical operations

- All operations are carried out by sending **HTTP requests**:
- Create a bucket/object: **POST**, along with the URI
- Update a bucket/object: **PUT/PATCH**, along with the URI
- Listing objects: **GET** on a bucket name
- Reading an object: **GET** on a full URI
- Removing an object: **Delete** on a bucket ID


Centralized system architectures

Basic Client–Server Model

Characteristics:

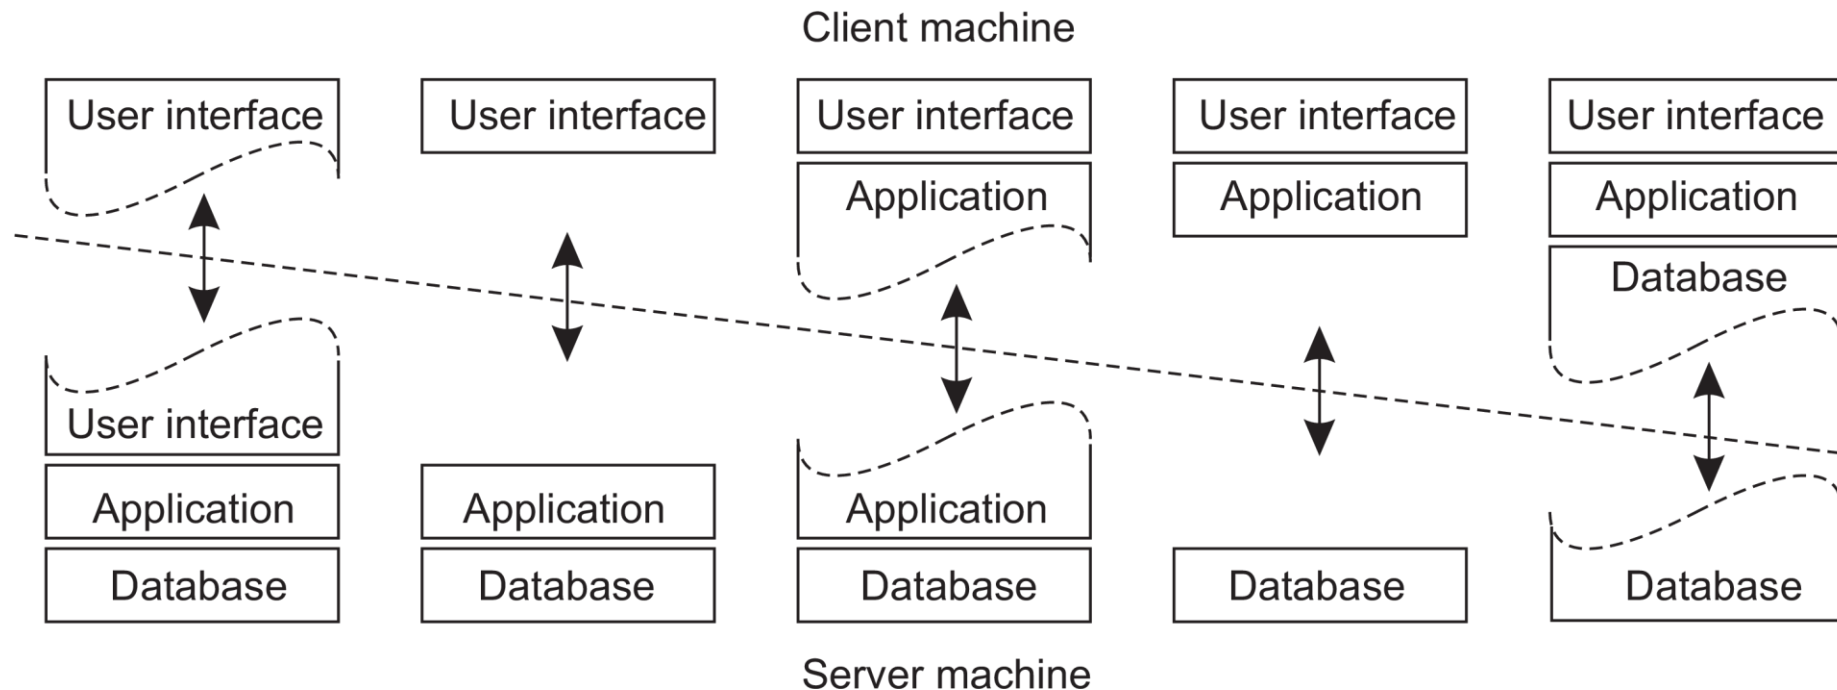
- There are processes offering services (**servers**)
- There are processes that use services (**clients**)
- Clients and servers can be on different machines
- Clients follow request/reply model with respect to using services

Centralized system architectures



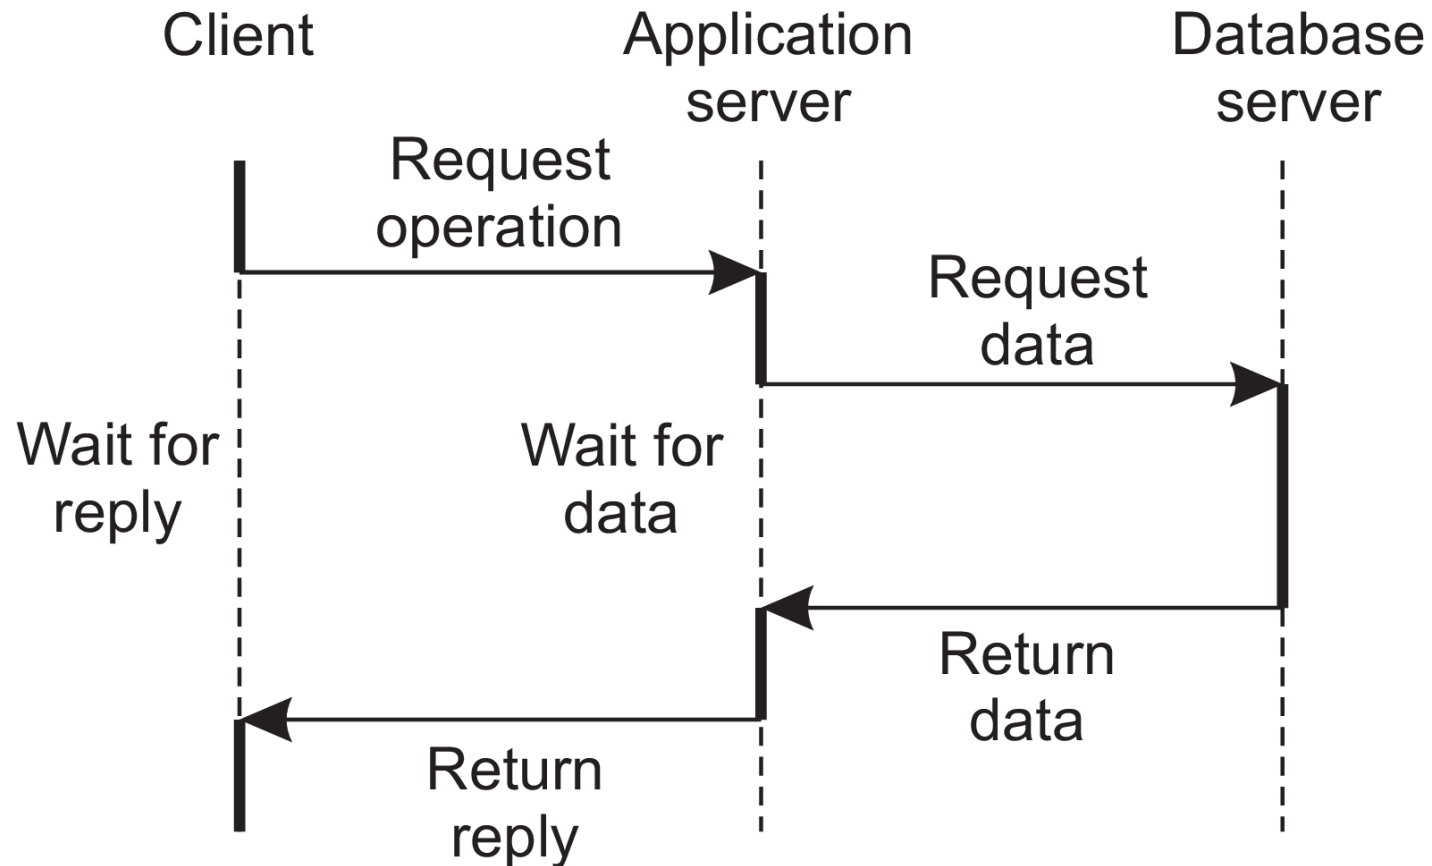
Multi-tiered centralized system architectures

- **Single-tiered**: dumb terminal/mainframe configuration
- **Two-tiered**: client/single server configuration
- **Three-tiered**: each layer on separate machine



Being client and server at the same time

- Three-tiered architecture




Skype's principle operation: A wants to contact B

Both A and B are on the public Internet

- A TCP connection is set up between A and B for control packets.
- The actual call takes place using UDP packets between negotiated ports

Edge-server architecture

- Systems deployed on the Internet where servers are placed at the edge of the network:
- The boundary between enterprise networks and the actual Internet.

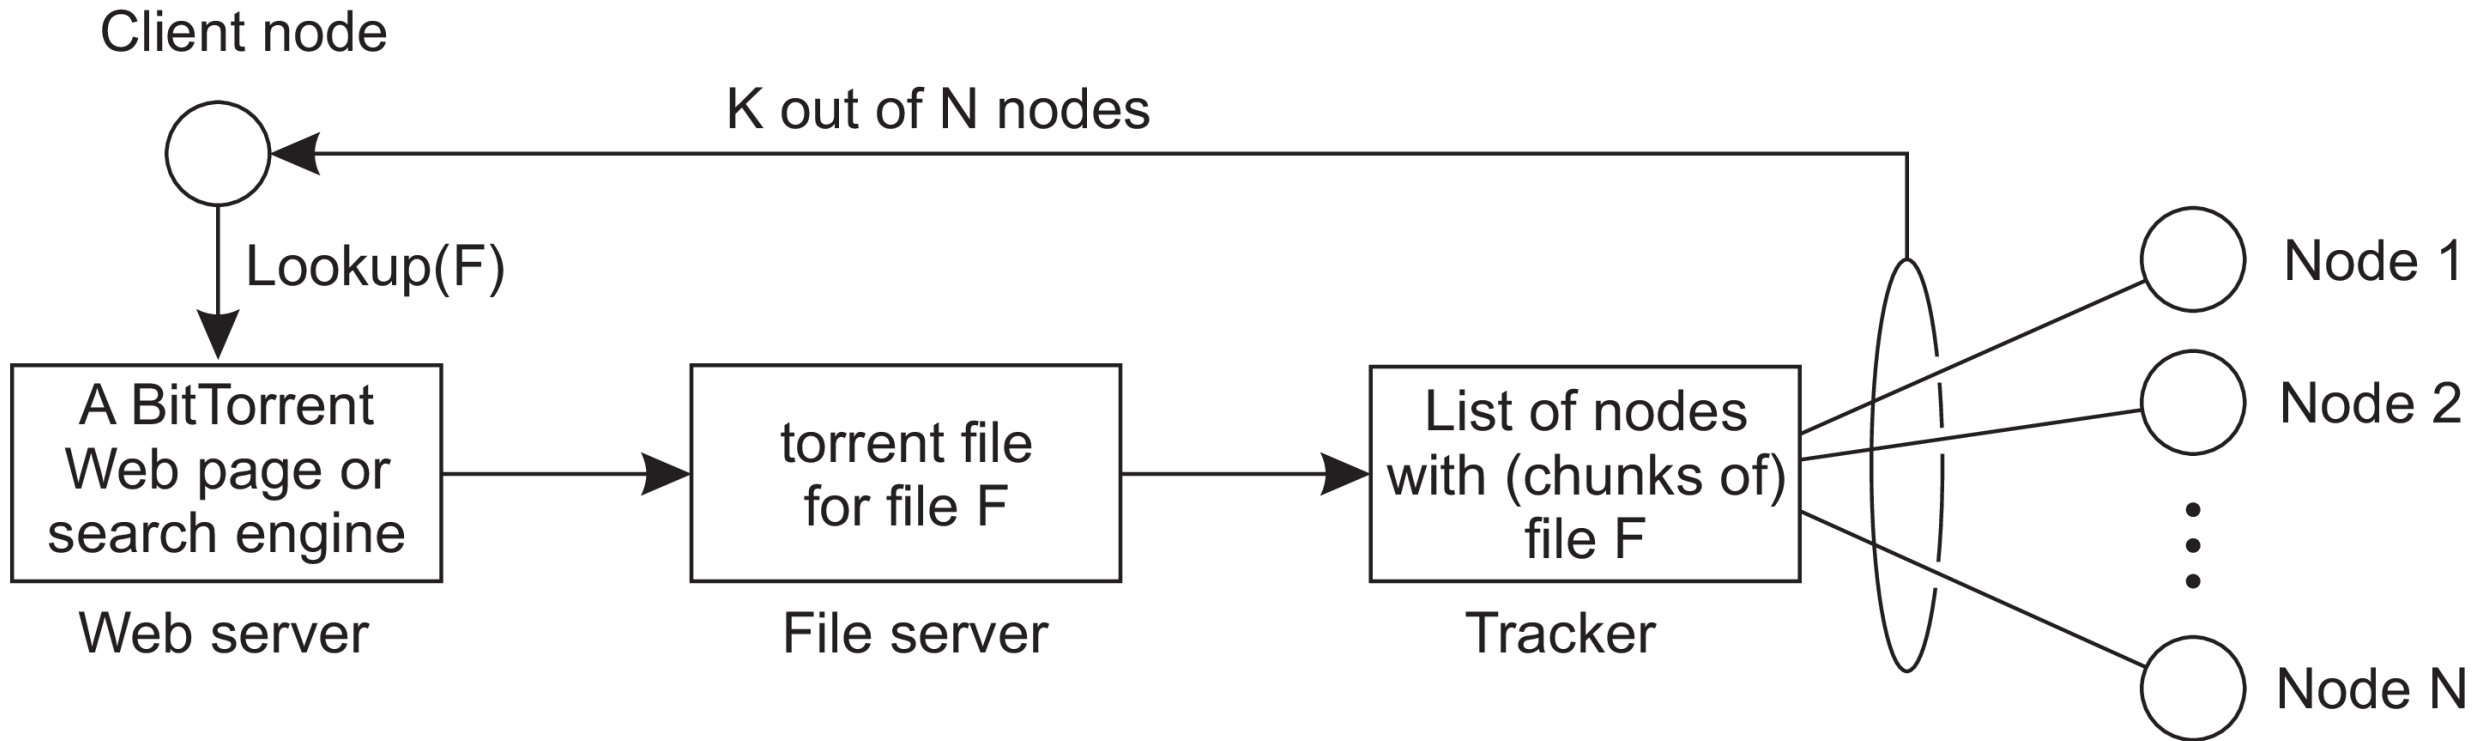


Collaboration: The BitTorrent case

Principle: search for a file F

- Lookup file at a global directory \Rightarrow returns a **torrent file**
- Torrent file contains reference to **tracker**: a server keeping an accurate account of **active** nodes that have (chunks of) F.
- P can join swarm, get a chunk for free, and then trade a copy of that chunk for another one with a peer Q also in the swarm.

Collaboration: The BitTorrent case



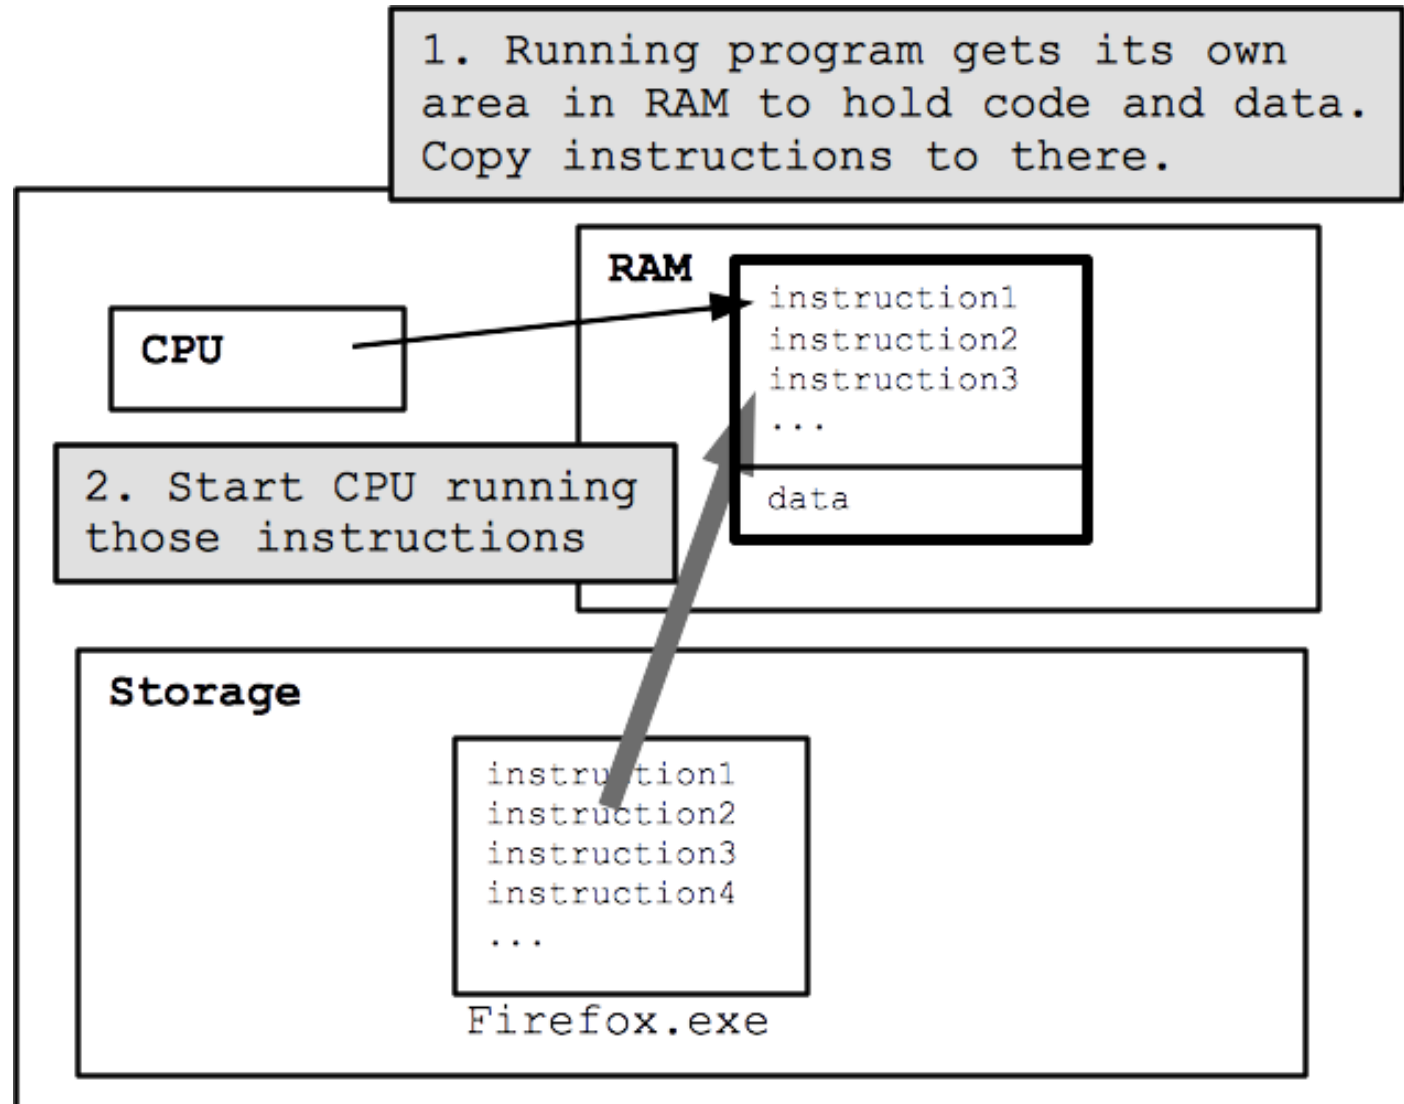
Introduction to threads

We build **virtual processors** in software, on top of physical processors:

Processor:

- Provides a set of instructions along with the capability of automatically executing a series of those instructions.

Processor





Introduction to threads

Thread:

- A minimal software processor in whose context a series of instructions can be executed. Saving a thread context implies stopping the current execution and saving all the data needed to continue the execution at a later stage.

Thread



threaded


Thread and Cores



4 MB Cache	6 MB Cache	8 MB Cache	16 MB Cache
2 Cores	4 Cores	4 Cores	8 Cores
4 Thread	8 Thread	8 Thread	16 Thread

What is the difference between a 'Thread' and a 'Core'?

- In hardware, a '**thread**' usually means a **logical core**. It may or may not be a **physical core**
- In software, a '**thread**' depends on the OS type, but is a single continuous piece of code executing



Introduction to threads

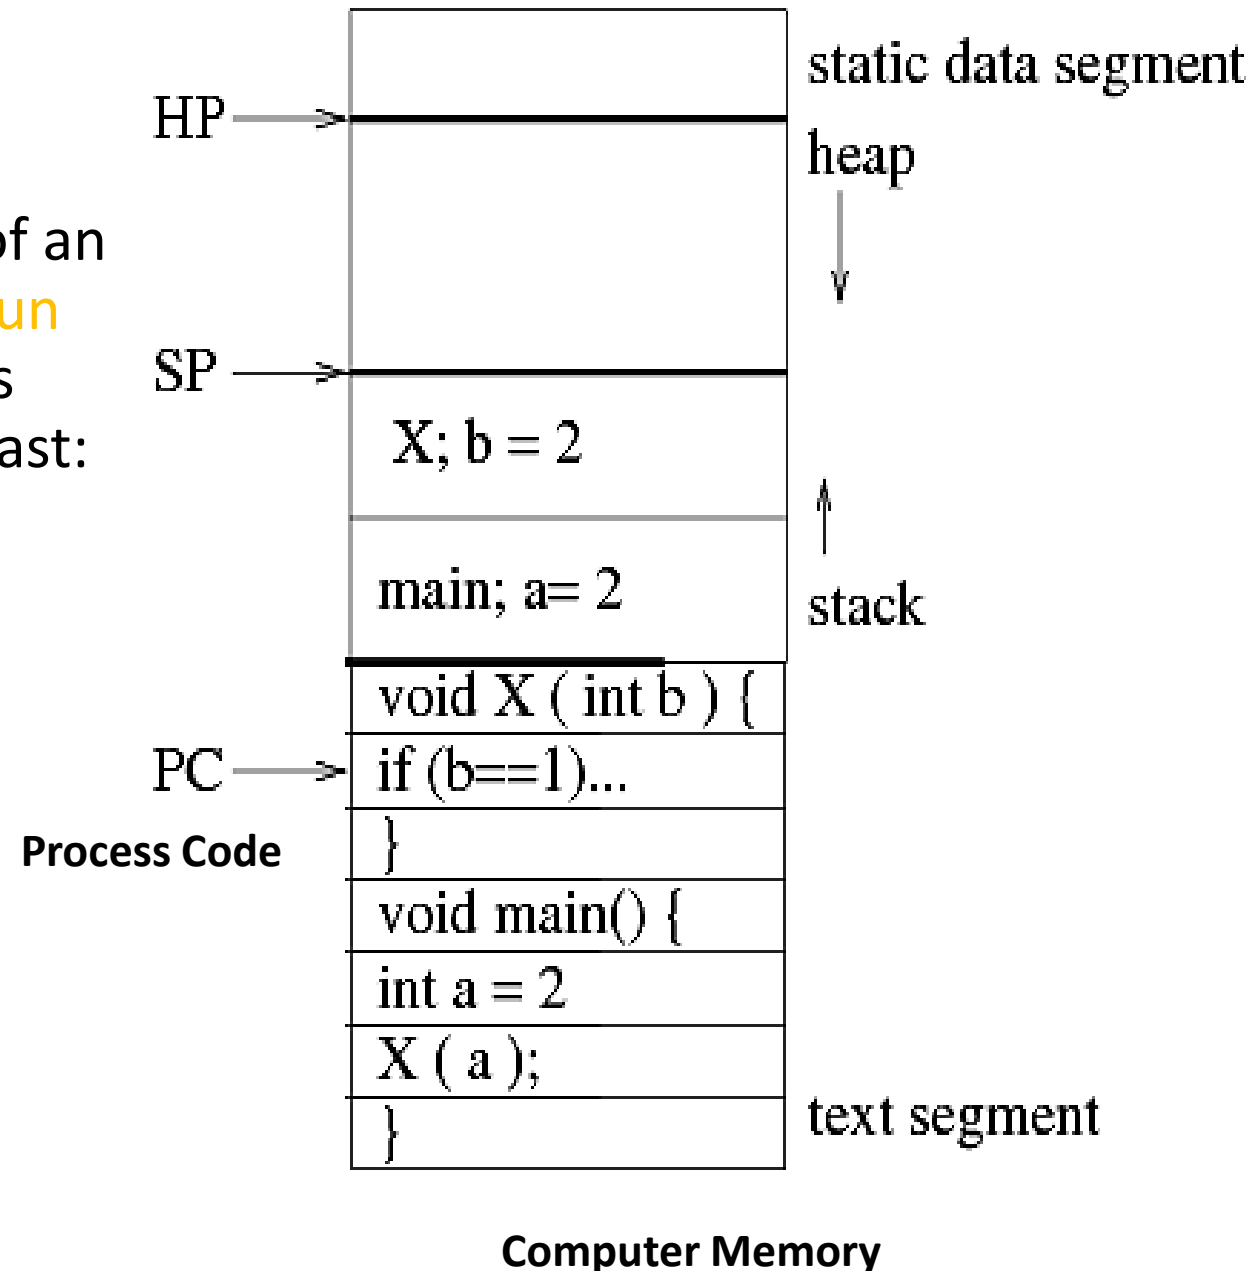
Process:

- A software processor in whose context one or more threads may be executed. Executing a thread, means executing a series of instructions in the context of that thread

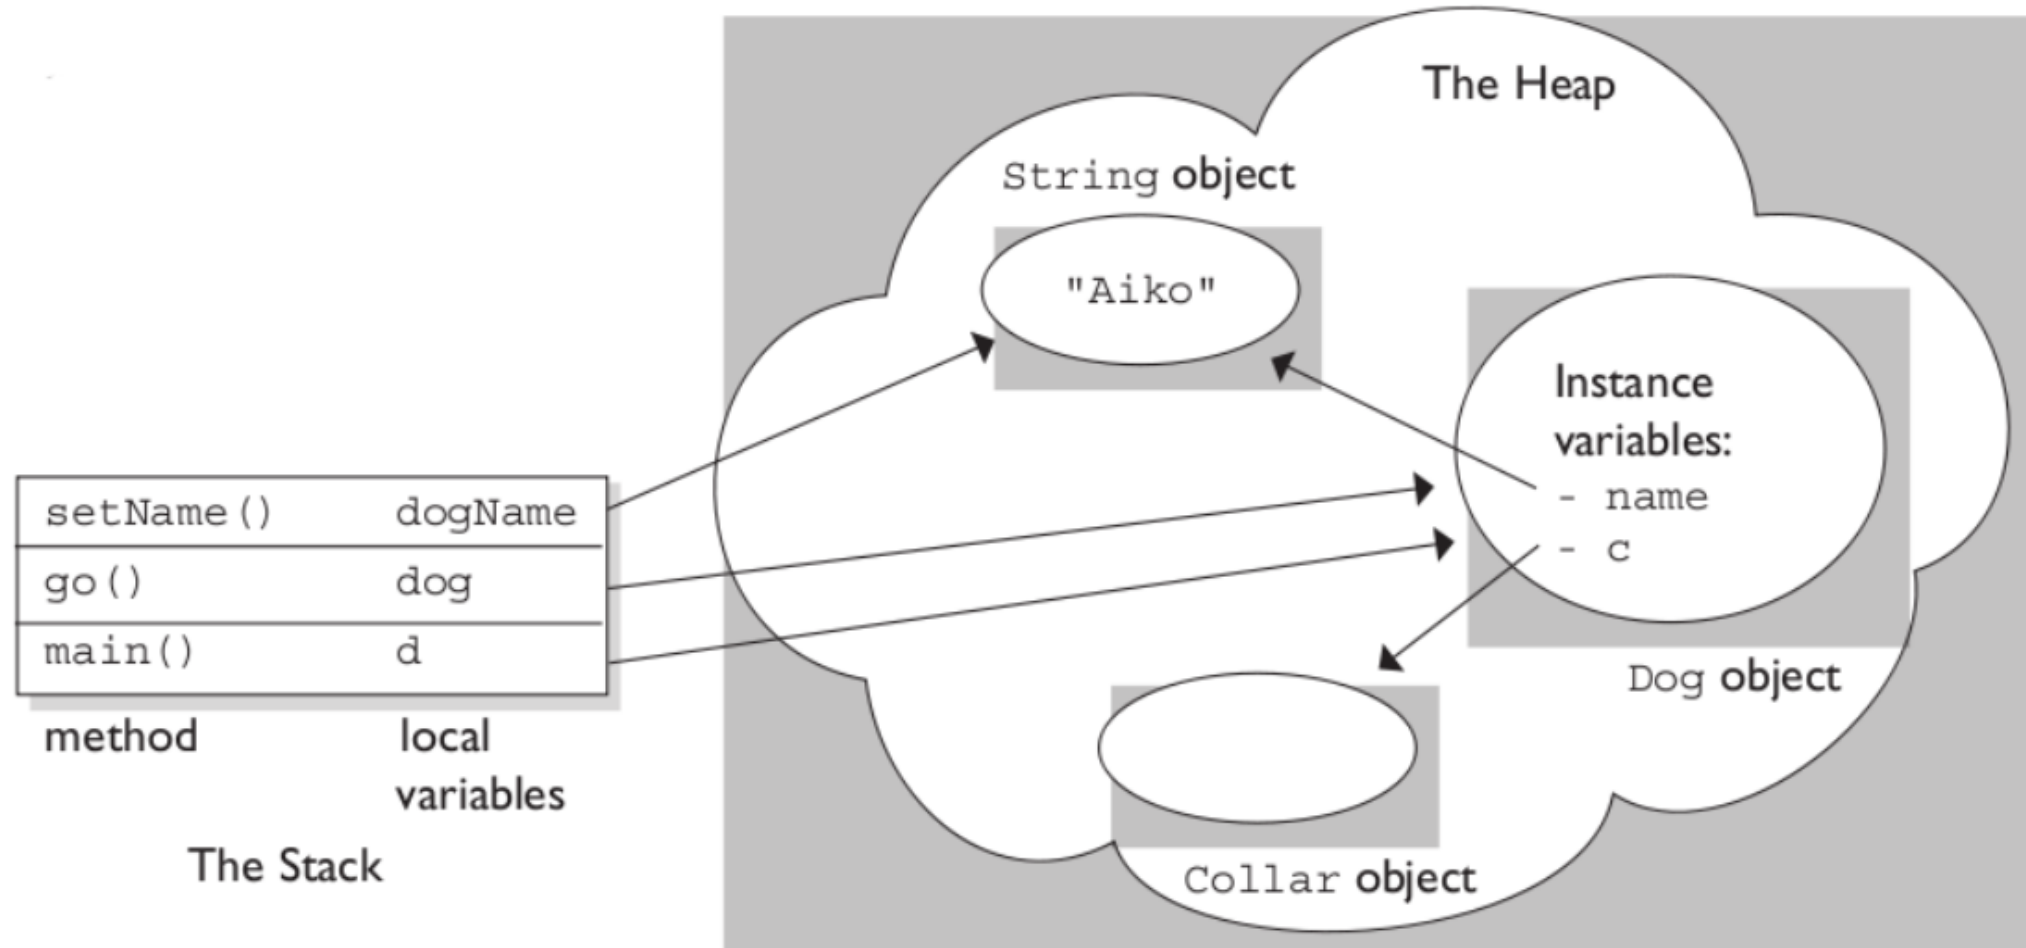
Process consists

A process is the **dynamic execution context** of an **executing program**. **Several processes may run concurrently**, each as a distinct entity with its own state. The process state consists of at least:

- the code for the running program
- the static data for running program
- a space in the heap for dynamic data
- the address of next instruction
- execution stack
- values of the CPU registers
- OS resources (e.g. open files)
- process execution state



Stack Vs Heap



Context switching

- **Processor context:** The minimal collection of values stored in the registers of a processor used for the execution of a series of instructions (e.g., stack pointer, addressing registers, program counter).
- **Thread context:** The minimal collection of values stored in registers and memory, used for the execution of a series of instructions (i.e., processor context, state).
- **Process context:** The minimal collection of values stored in registers and memory, used for the execution of a thread (i.e., thread context).


Context switching

- Threads share the same address space. Thread context switching can be done entirely independent of the operating system.
- Process switching is generally (somewhat) more expensive as it involves getting the OS in the loop, i.e., trapping to the kernel.
- Creating and destroying threads is much cheaper than doing so for processes.

Why use threads

- **Avoid needless blocking:** a single-threaded process will block when doing I/O; in a multi-threaded process, the operating system can switch the CPU to another thread in that process.
- **Exploit parallelism:** the threads in a multi-threaded process can be scheduled to run in parallel on a multiprocessor or multicore processor.
- **Avoid process switching:** structure large applications not as a collection of processes, but through multiple threads.

Avoid process switching



Thread context switching may be faster than process context switching

Using threads at the client side

Hiding network latencies:

- **Web browser** scans an incoming HTML page, and finds that **more files need to be fetched.**
- **Each file is fetched by a separate thread,** each doing a (blocking) HTTP request.
- As files come in, the browser displays them

Using threads at the server side

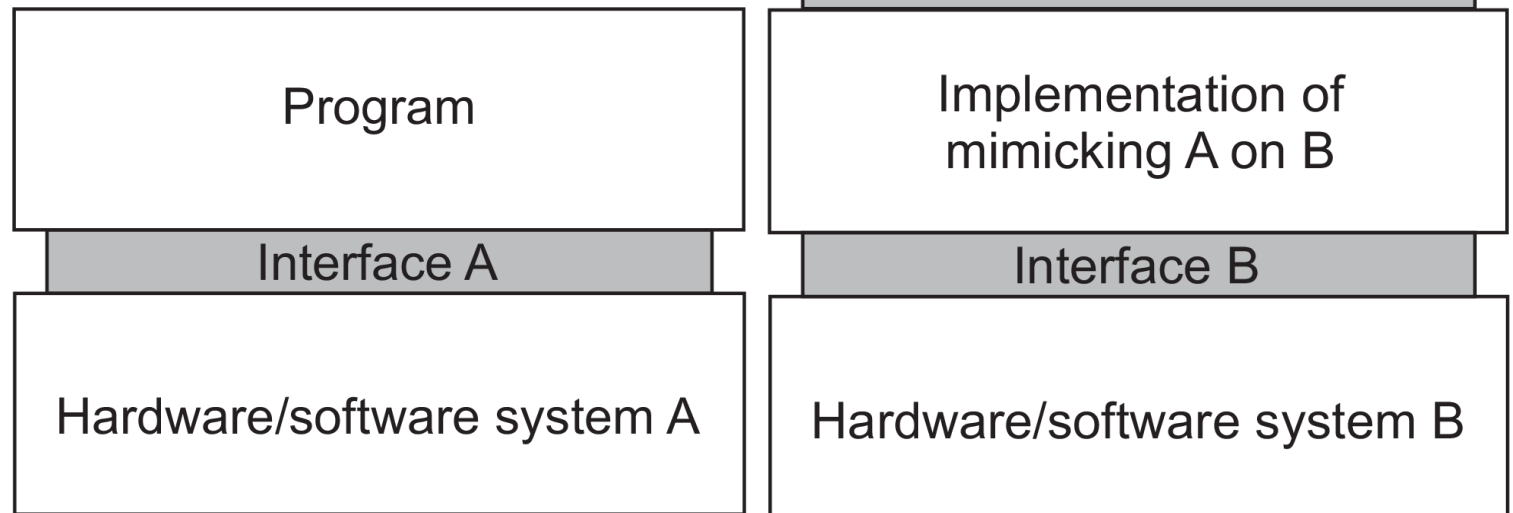
Improve performance

- Starting a thread is cheaper than starting a new process.
- Having a single-threaded server prohibits simple scale-up to a multiprocessor system.
- As with clients: hide network latency by reacting to next request while previous one is being replied.

Virtualization

Virtualization is important:

- Hardware **changes faster** than software
- Ease of **portability** and code migration
- **Isolation** of failing or attacked components



VMs and cloud computing

Three types of cloud services

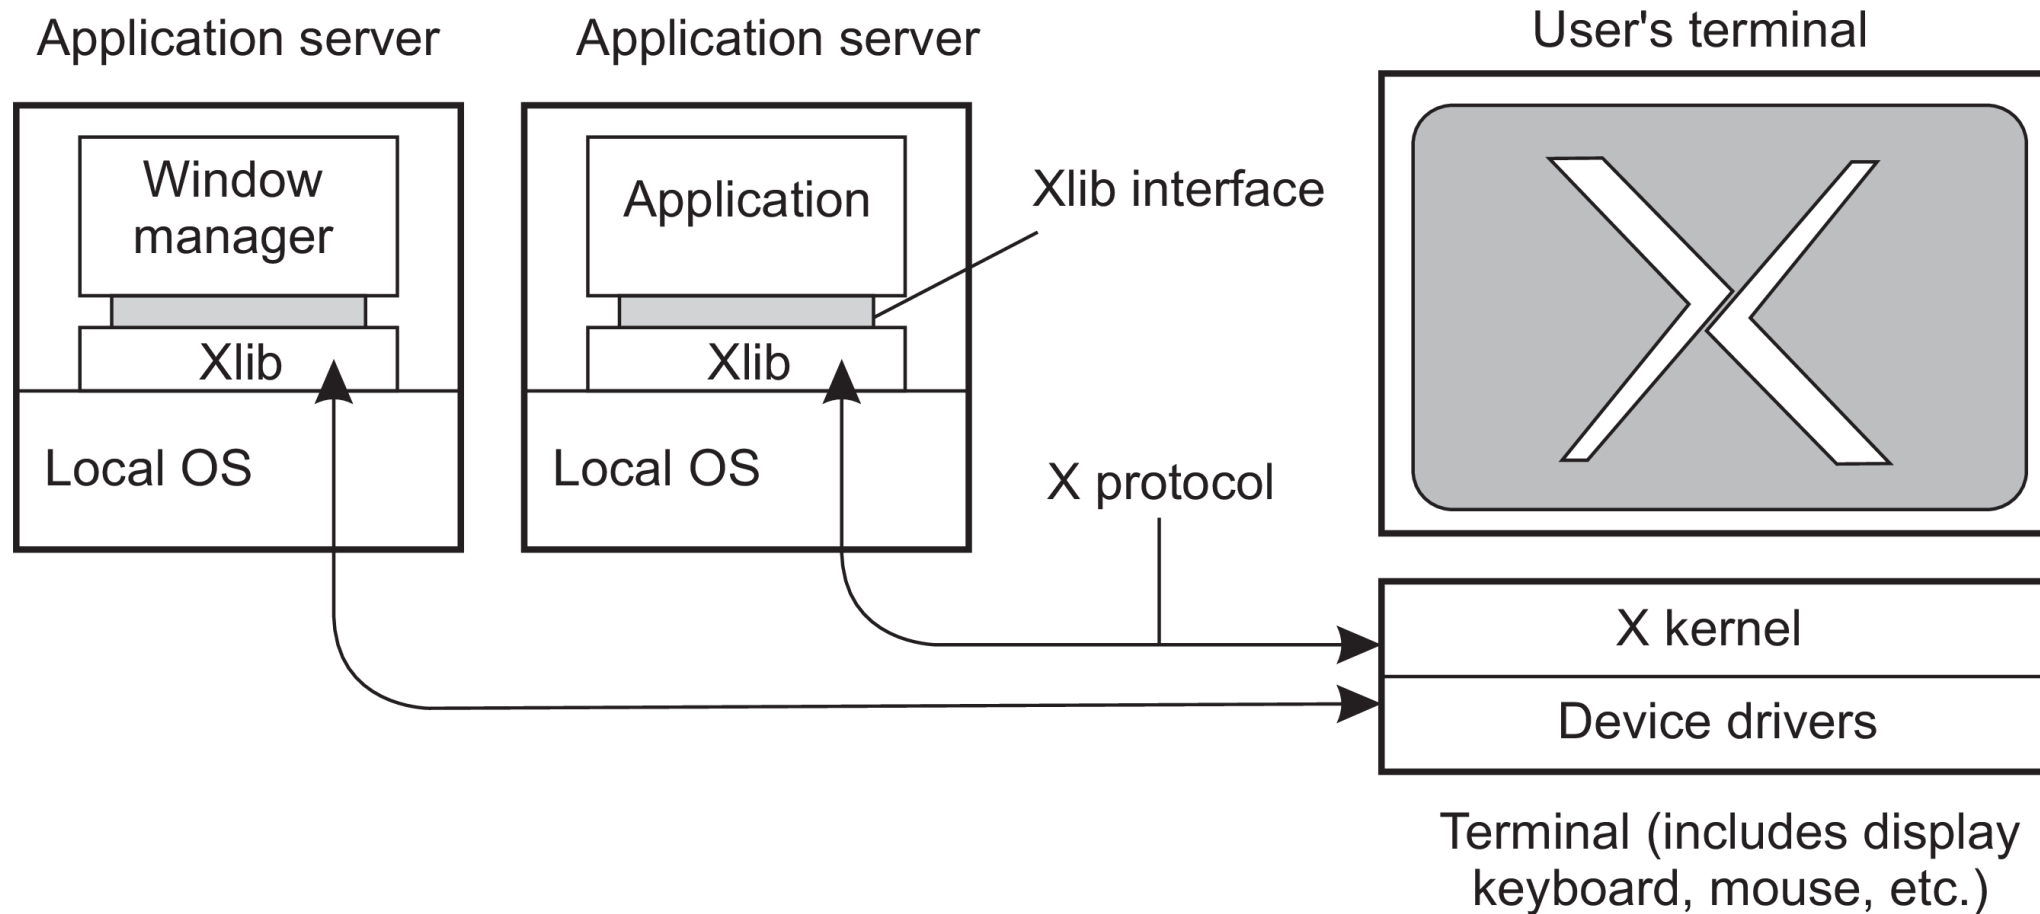
- **Infrastructure-as-a-Service** covering the basic infrastructure
- **Platform-as-a-Service** covering system-level services
- **Software-as-a-Service** containing actual applications

IaaS

Instead of renting out a physical machine, a cloud provider will rent out a VM


Almost complete isolation between customers

Example X Windows



Client-side software

- Access transparency: client-side stubs for RPCs
- Location/migration transparency: let client-side software keep track of actual location
- Replication transparency: multiple invocations handled by client stub:



Servers

- A process implementing a **specific service** on behalf of a **collection of clients**.
- It waits for an **incoming request from a client** and subsequently ensures that the request is taken care of,
- After which it waits for the **next incoming request**.

Concurrent servers

- **Iterative server:**

Server handles the request before attending a next request.

- **Concurrent server:**

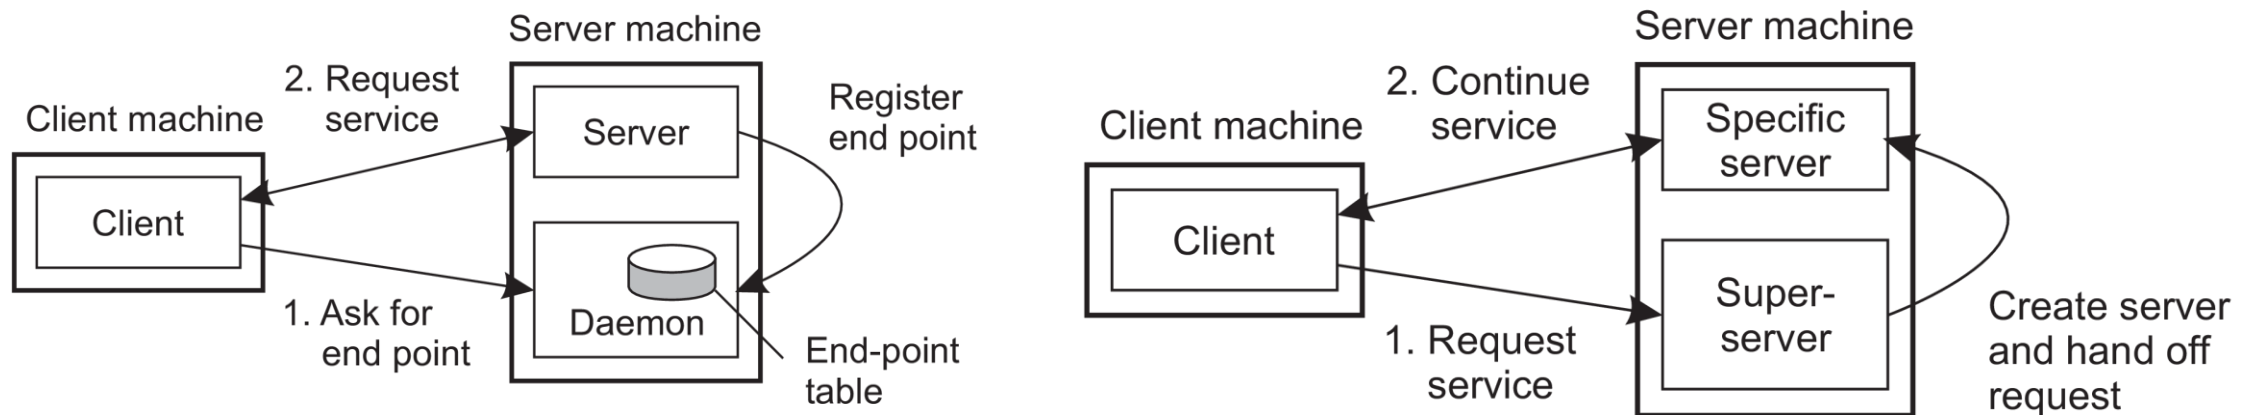
Uses a **dispatcher**, which picks up an incoming request that is then passed on to a separate thread/process.

Concurrent servers: they can easily **handle multiple requests**, notably in the presence of **blocking operations**

Contacting a server

- Observation: most services are tied to a specific port

ftp-data	20	File Transfer [Default Data]
ftp	21	File Transfer [Control]
telnet	23	Telnet
Smtp	25	Simple Mail Transfer
www	80	Web (HTTP)



Out-of-band communication

Issue

- Is it possible to **interrupt** a server once it has accepted (or is in the process of accepting) a service request?

Solution 1: Use a separate port for urgent data

- Server has a separate thread/process for urgent messages
- Urgent message comes in \Rightarrow **associated request is put on hold**
- Note: we require **OS supports priority-based scheduling**

Solution 2: Use facilities of the transport layer

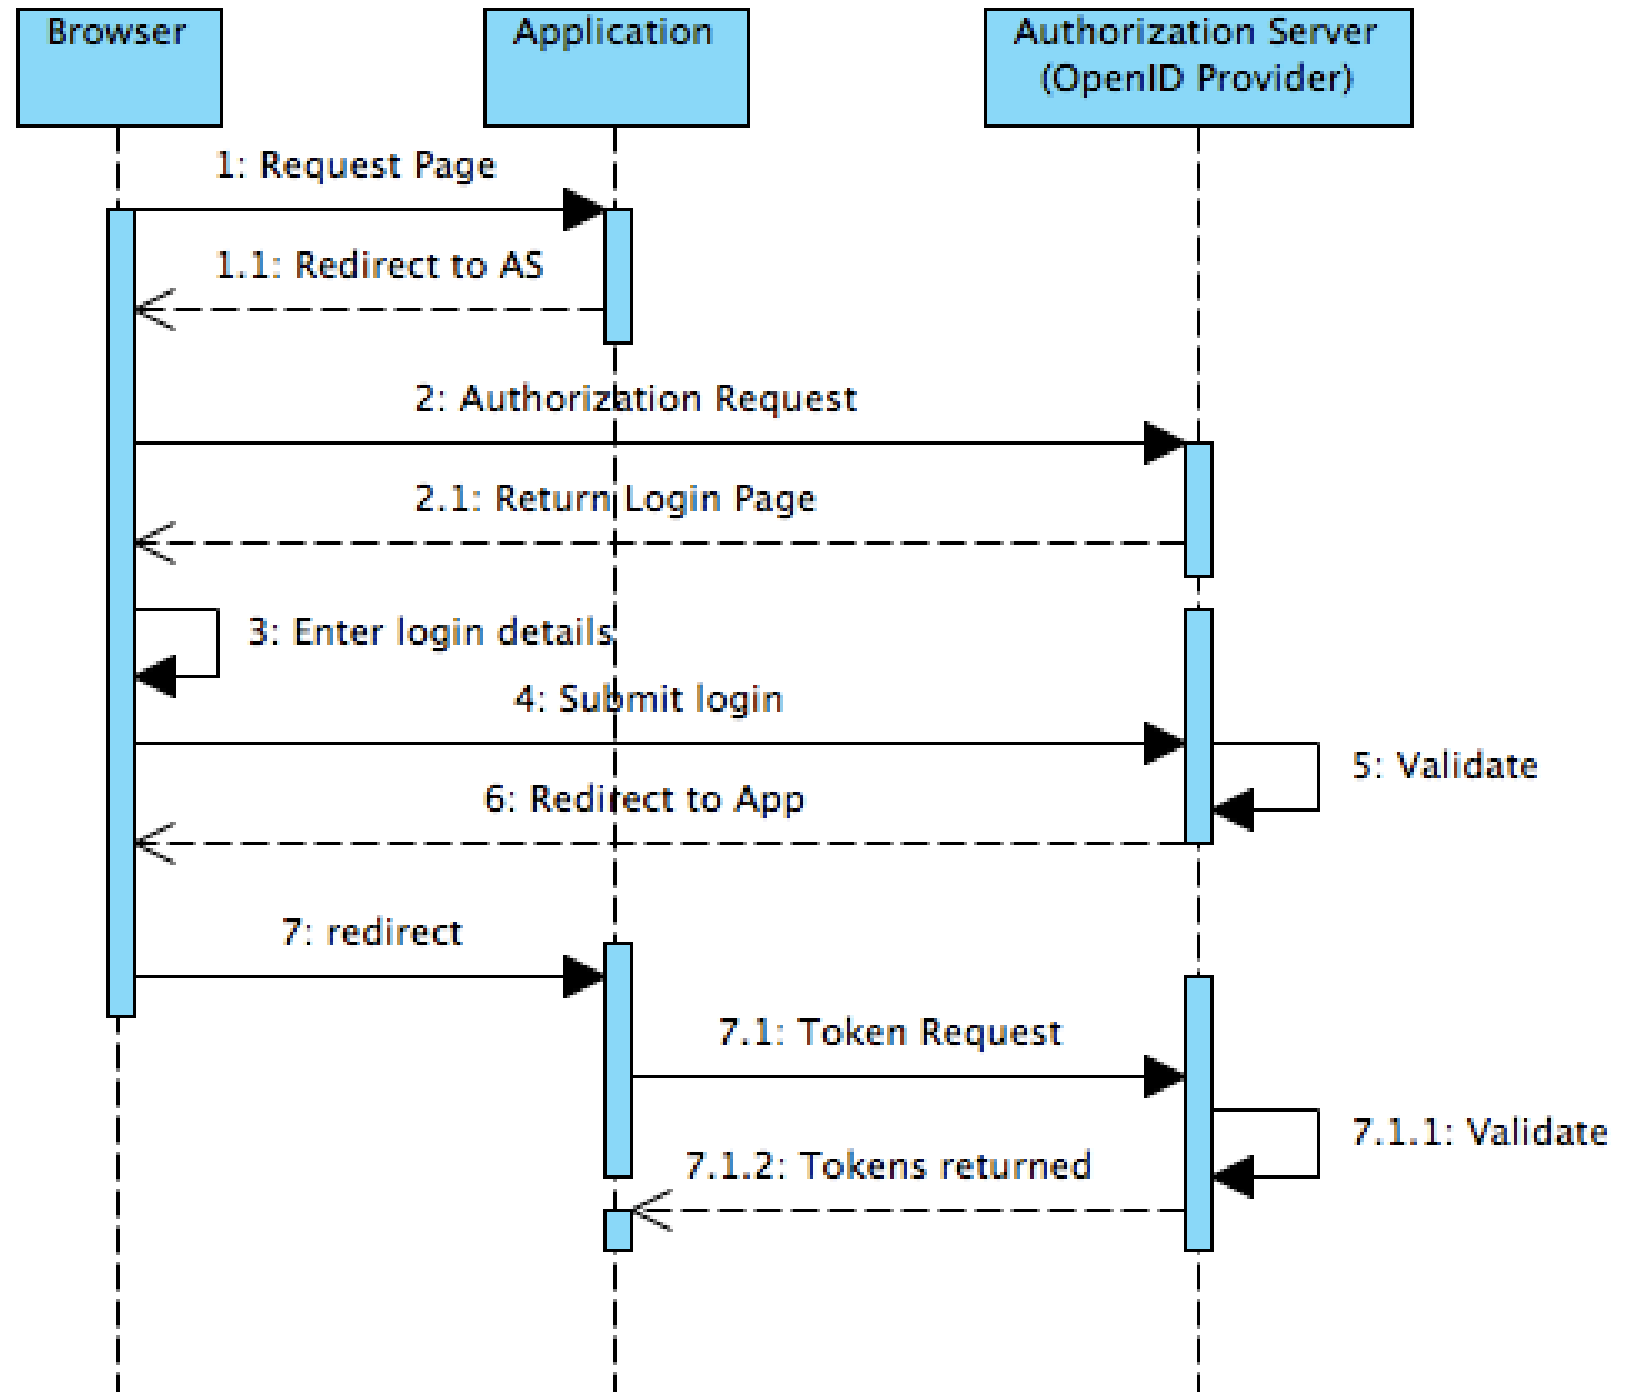
- Example: TCP allows for urgent messages in same connection
- Urgent messages can be caught using OS signaling techniques

Servers and state

Stateless servers

- Never keep **accurate** information about the status of a client after having handled a request:
- Don't record whether a file has been opened (simply close it again after access)
- Don't promise to invalidate a client's cache
- Don't keep track of your clients

Client Request Server



Servers and state

Consequences

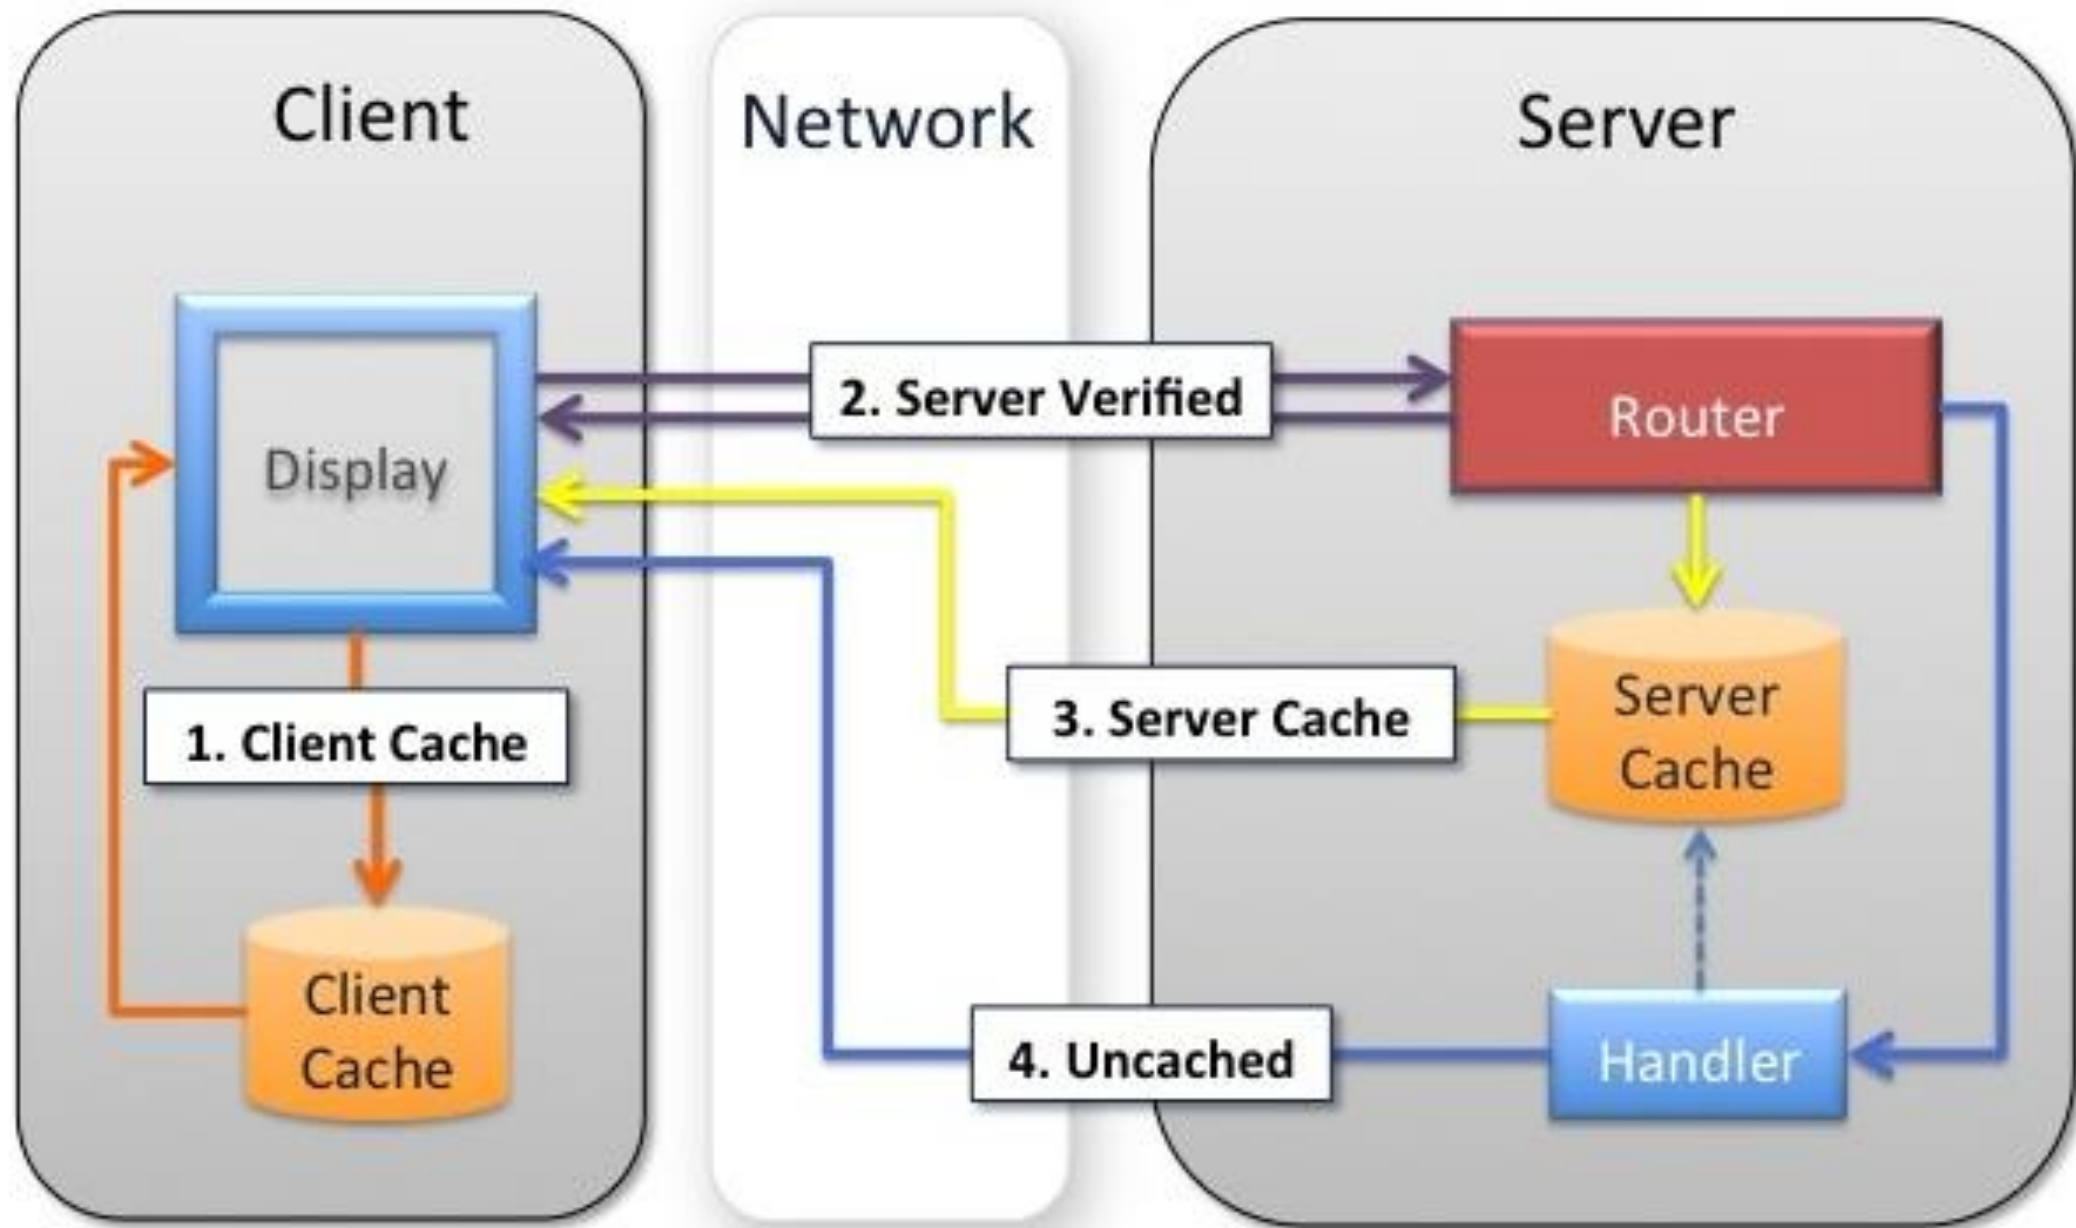
- Clients and servers are **completely independent**
- **State inconsistencies** due to client or server crashes are reduced
- Possible **loss of performance** because, e.g., a server cannot anticipate client behavior (think of prefetching file blocks)

Servers and state

Statefull servers

- Keeps track of the status of its clients:
- Record that a file has been **opened**, so that prefetching can be done
- Knows which **data a client has cached**, and allows clients to keep local copies of shared **data**

Appweb Caching Options




Servers and state

Observation

- The performance of statefull servers can be **extremely high**,
- Provided clients are allowed to keep **local copies**.
- As it turns out, **reliability** is often not a major problem.

Three different tiers



The first tier is generally responsible for passing requests to an appropriate server: **request dispatching**

Request Handling

- Having the first tier handle all communication from/to the cluster may lead to a **bottleneck**: A solution: TCP handoff

