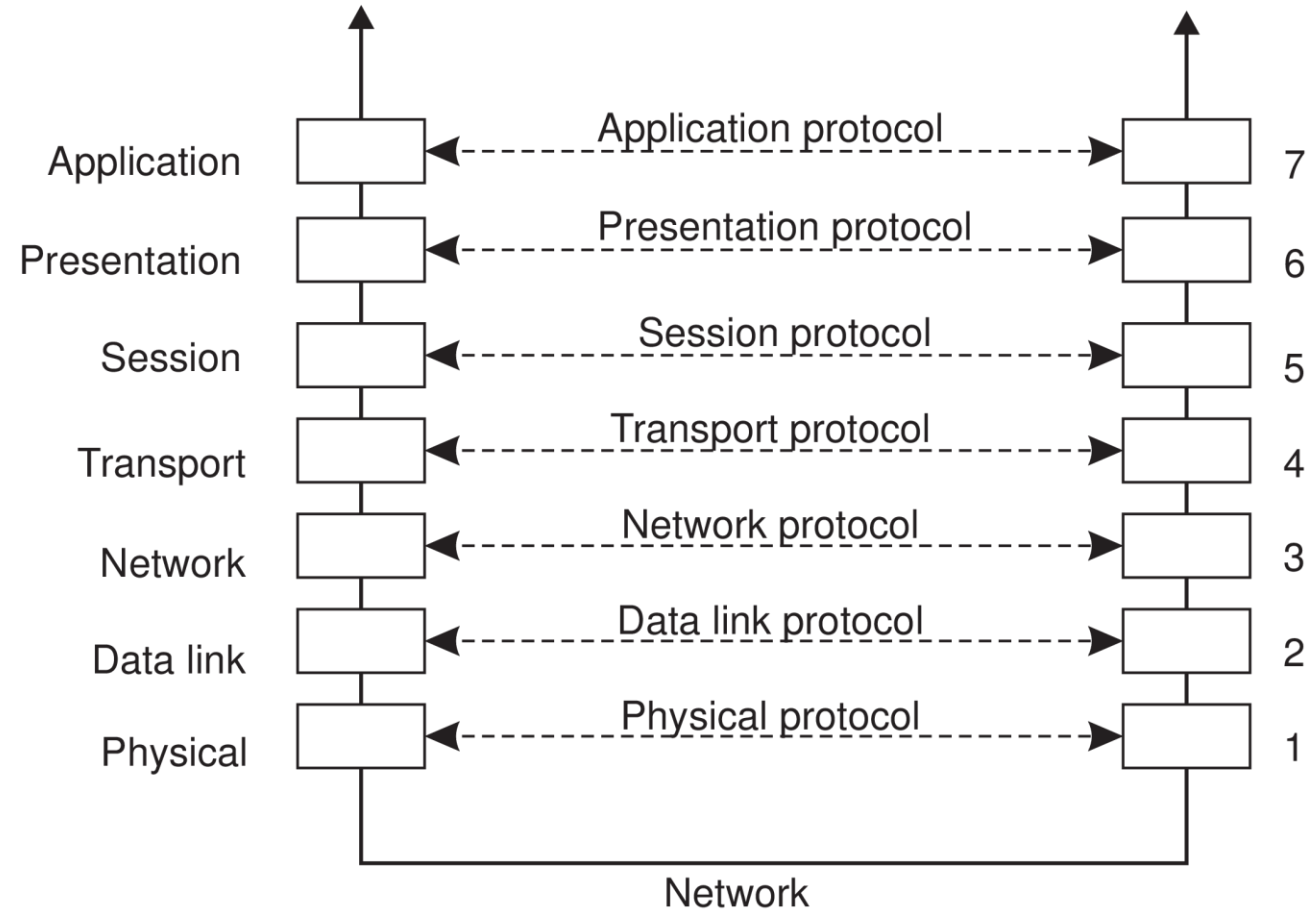


Distributed Systems

Basic networking model



Low-level layers

- **Physical layer:** contains the specification and implementation of bits, and their transmission between sender and receiver
- **Data link layer:** prescribes the transmission of a series of bits into a frame to allow for error and flow control
- **Network layer:** describes how packets in a network of computers are to be **routed**.

Transport Layer

- The **transport layer** provides the actual communication facilities for most distributed systems.

Standard Internet protocols:

- **TCP**: connection-oriented, reliable, stream-oriented communication
- **UDP**: unreliable (best-effort) datagram communication

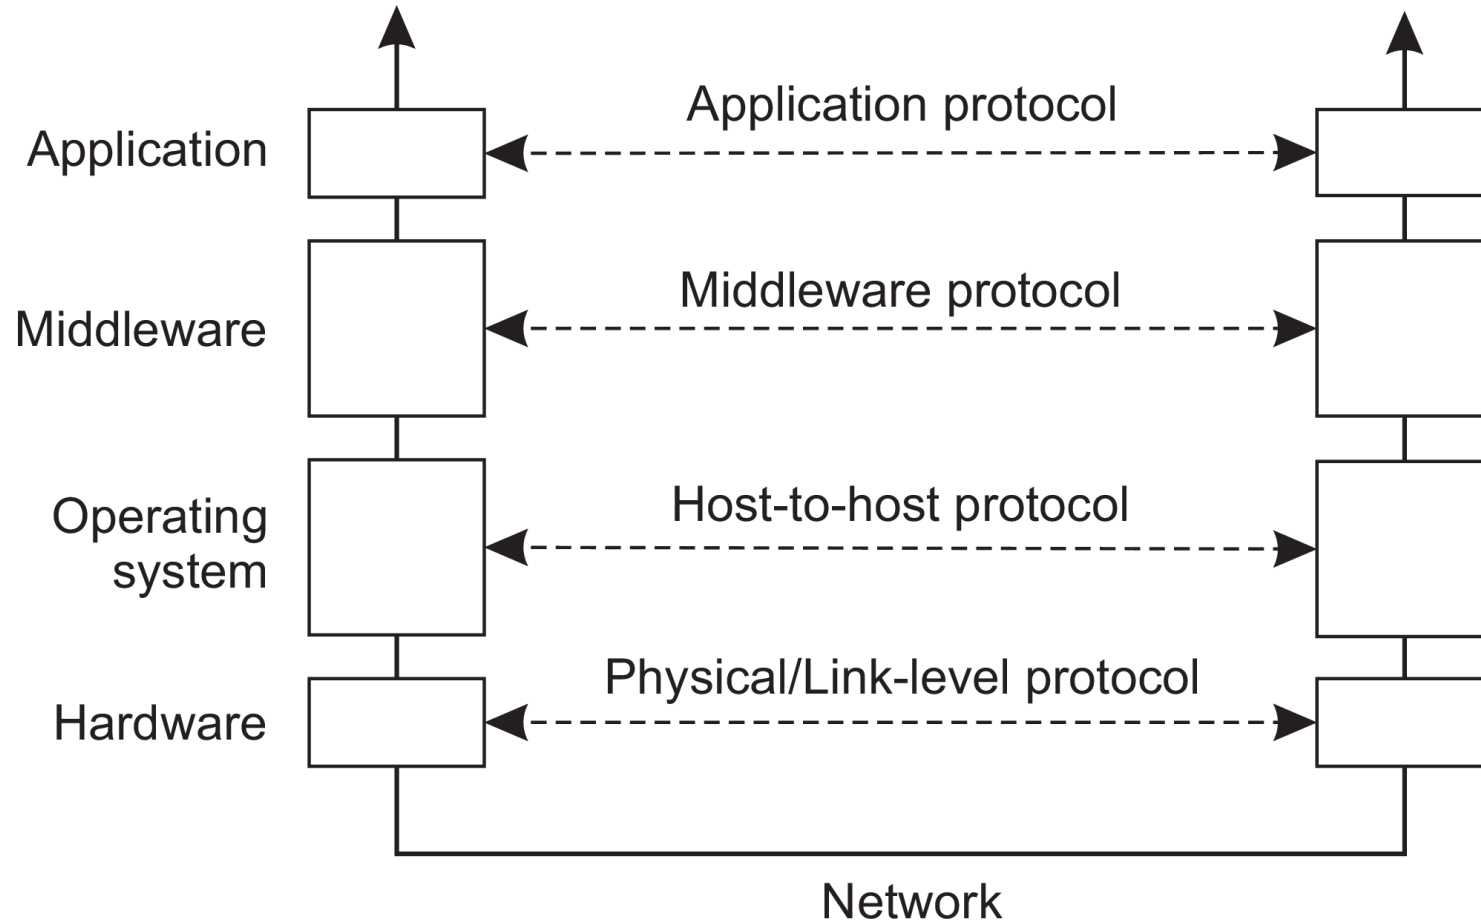
Middleware layer

- Middleware is invented to provide **common** services and protocols that can be used by many **different** applications
- A rich set of **communication protocols**
- **(Un)marshaling** of data, necessary for integrated systems
- **Naming protocols**, to allow easy sharing of resources
- **Security protocols** for secure communication
- **Scaling mechanisms**, such as for replication and caching


Note

- What remains are truly **application-specific** protocols... **such as?**

An adapted layering scheme



Types of communication



Types of communication

- **Transient** versus **persistent** communication
- **Asynchronous** versus **synchronous** communication

Transient communication:

- Comm. server discards message when it cannot be delivered at the next server, or at the receiver.

Persistent communication:

- A message is stored at a communication server as long as it takes to deliver it.

At **request submission**

At **request delivery**

After **request processing**

Client/Server

- Client/Server computing is generally based on a model of **transient synchronous communication**:
- Client and server have to be active at time of **communication**
- Client issues request and **blocks** until it receives reply
- Server essentially waits only for **incoming requests**, and subsequently processes them

Drawbacks synchronous communication

- **Client** cannot do any other **work** while waiting for reply
- **Failures** have to be handled immediately: the client is waiting
- **The model** may simply not be appropriate (mail, news)

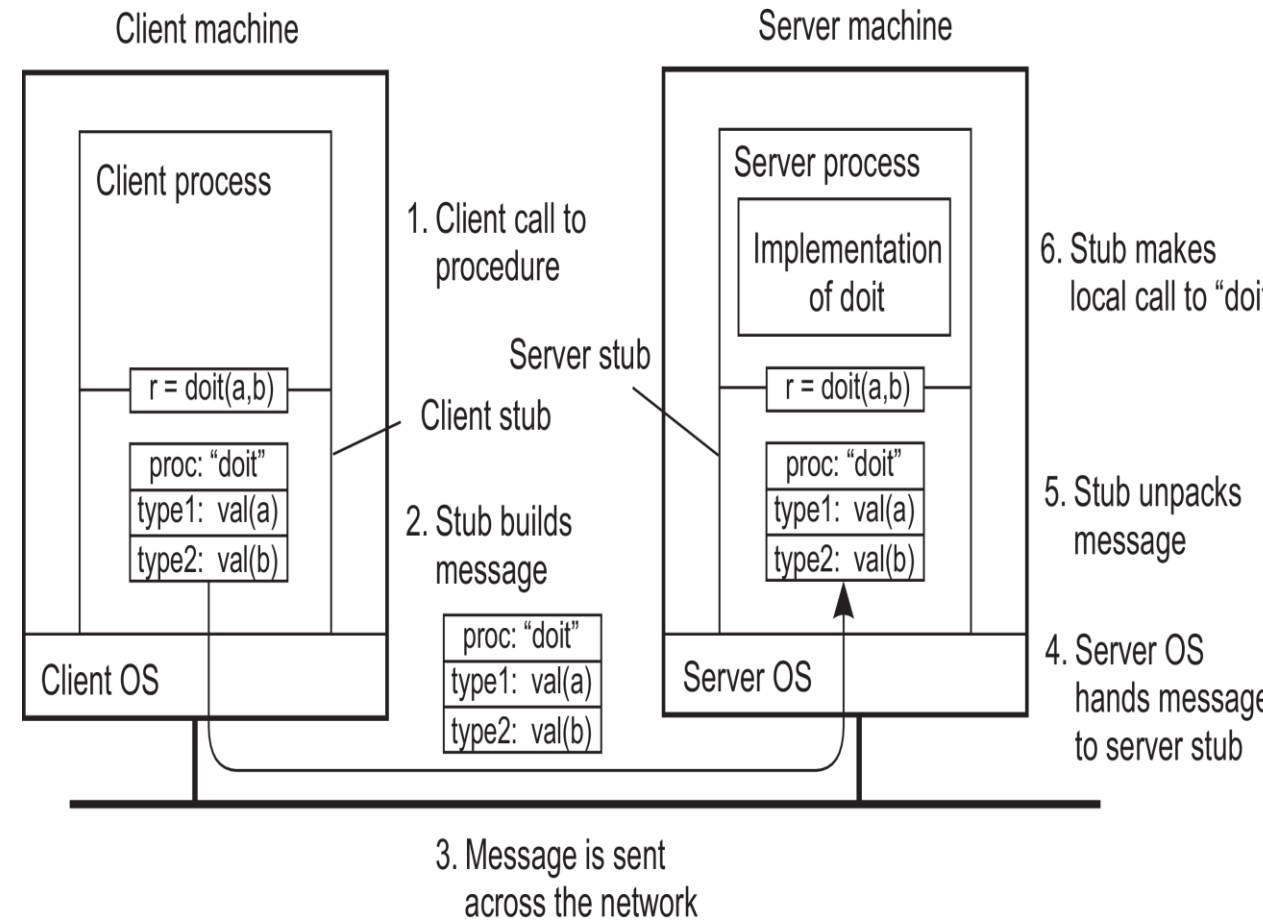
Messaging

Message-oriented middleware

- Aims at high-level **persistent asynchronous communication**:
- Processes send each other **messages**, which are **queued**
- **Sender** need not wait for **immediate reply**, but can do other things
- **Middleware** often ensures fault tolerance

Basic RPC operation

1. Client procedure calls client stub.
2. Stub builds message; calls local OS.
3. OS sends message to remote OS.
4. Remote OS gives message to stub.
5. Stub unpacks parameters; calls server.
6. Server does local call; returns result to stub.
7. Stub builds message; calls OS.
8. OS sends message to client's OS.
9. Client's OS gives message to stub.
10. Client stub unpacks result; returns to client.




RPC: Parameter passing

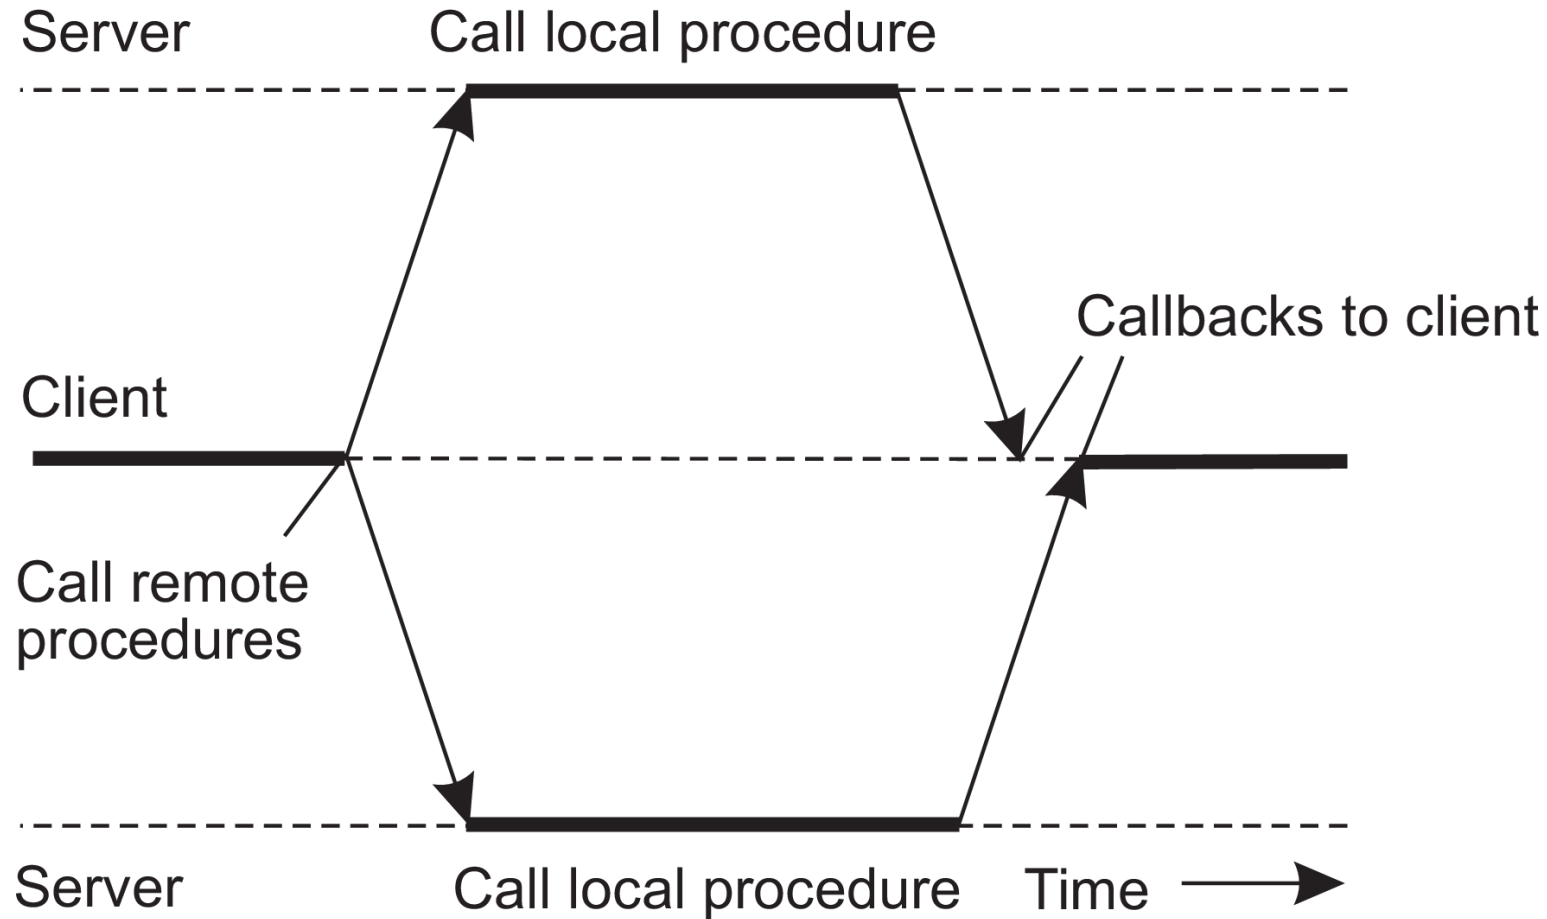
There's more than just wrapping parameters into a message

- Client and server machines may have different data representations (think of byte ordering)
- Wrapping a parameter means transforming a value into a sequence of bytes
- Client and server have to agree on the same encoding:
- How are basic data values represented (integers, floats, characters)
- How are complex data values represented (arrays, unions)
- * Client and server need to properly interpret messages, transforming them into machine-dependent representations.

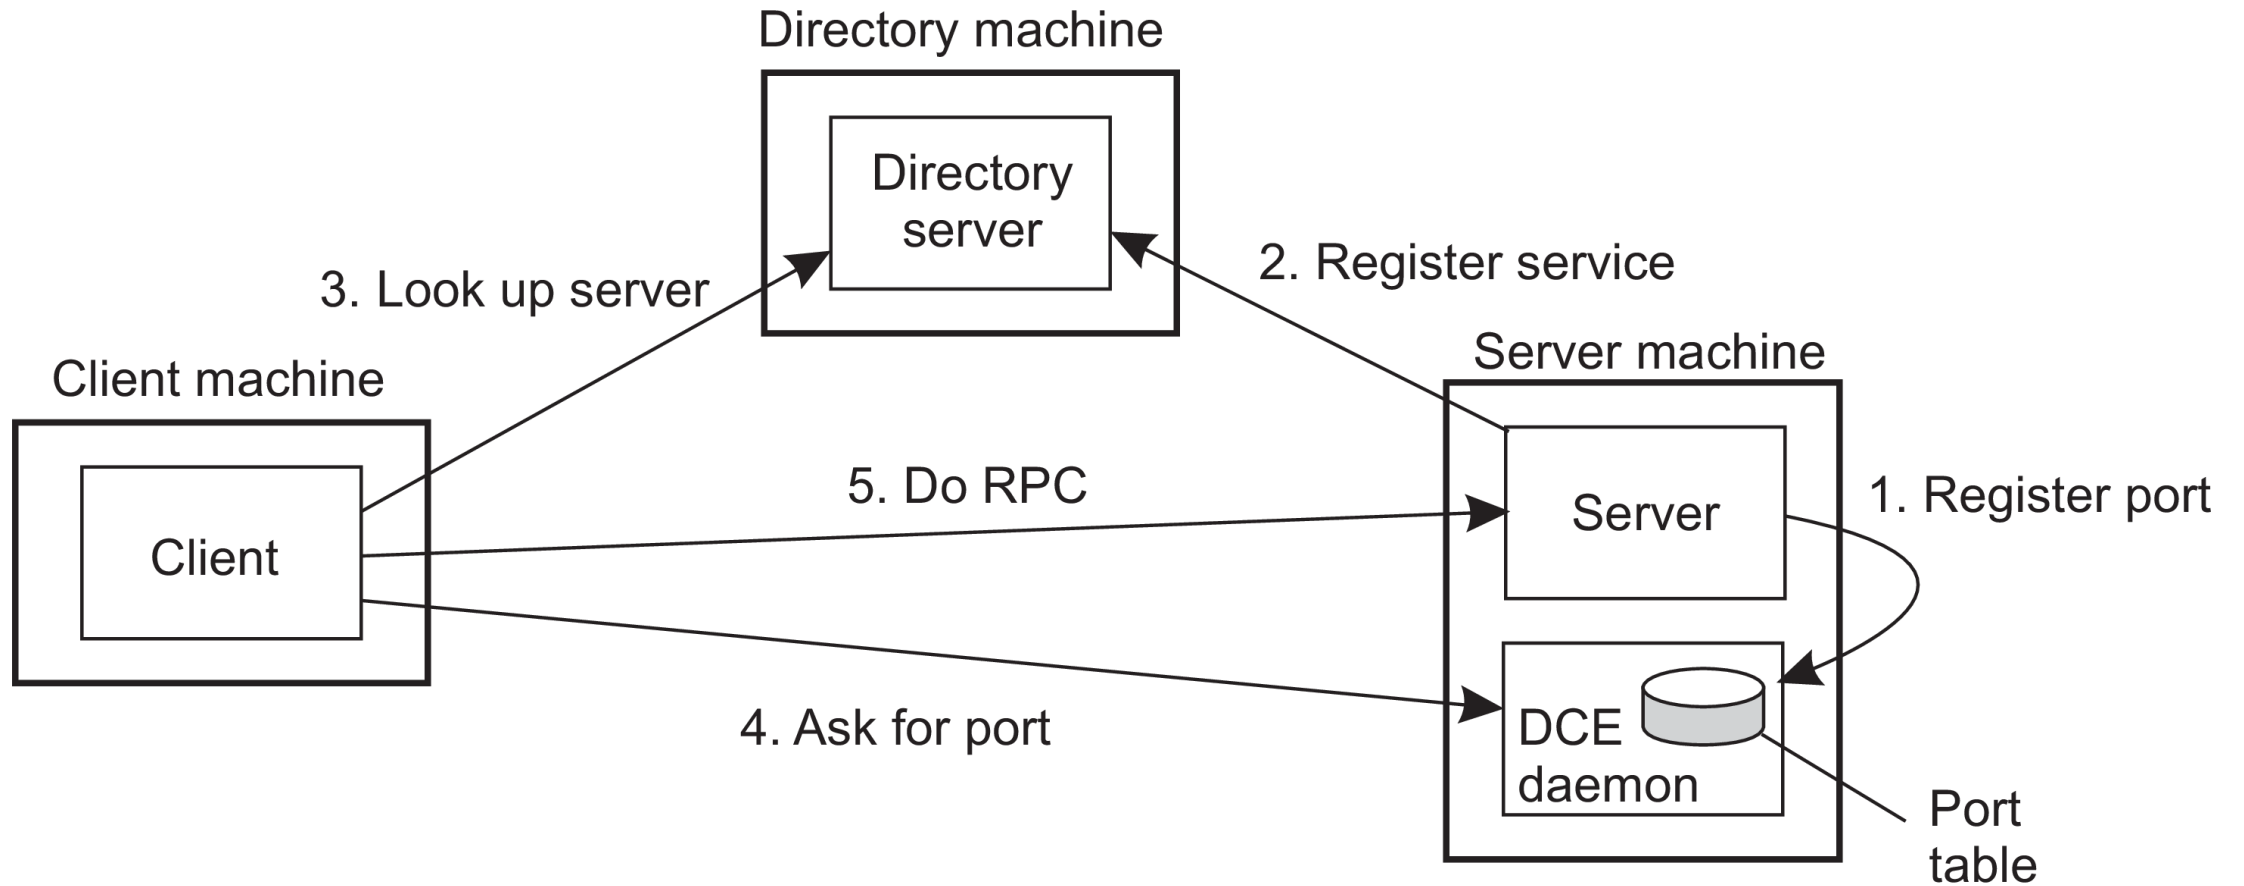
Asynchronous RPCs



Sending out multiple RPCs



Client-to-server binding




Message-oriented middleware

- Asynchronous persistent communication through support of middleware-level queues. **Queues** correspond to buffers at communication servers.

Operation	Description
put	Append a message to a specified queue
get	Block until the specified queue is nonempty, and remove the first message
Poll	Check a specified queue for messages, and remove the first. Never block
notify	Install a handler to be called when a message is put into the specified queue

General model Queue managers

- Queues are managed by **queue managers**. An application can put messages only into a **local** queue. Getting a message is possible by extracting it from a **local** queue only \Rightarrow queue managers need to **route** messages.



Message broker

- Message queuing systems assume a **common messaging protocol**: all applications agree on message format (i.e., structure and data representation)

Broker handles application heterogeneity in an MQ system

- Transforms incoming messages to target format
- Very often acts as an **application gateway**
- May provide **subject-based** routing capabilities (i.e., **publish-subscribe** capabilities)

Message broker: general architecture

